

# MANOJ KUMAR - 2048015

In [1]:

```
#Importing libraries  
import numpy as np  
import pandas as pd
```

In [2]:

```
#Importing the visualisation libraries  
import matplotlib.pyplot as plt  
import seaborn as sns  
%matplotlib inline
```

In [3]:

```
#Reading the data  
kidneyDisease = pd.read_csv('kidney_disease.csv')
```

## Exploratory Data Analysis

In [4]:

```
#Shape of the dataset  
kidneyDisease.shape
```

Out[4]:

```
(400, 26)
```

In [5]:

```
kidneyDisease.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    400 non-null    int64
1   age                  391 non-null    float64
2   bp                   388 non-null    float64
3   sg                   353 non-null    float64
4   al                   354 non-null    float64
5   su                   351 non-null    float64
6   rbc                  248 non-null    object
7   pc                   335 non-null    object
8   pcc                  396 non-null    object
9   ba                   396 non-null    object
10  bgr                  356 non-null    float64
11  bu                   381 non-null    float64
12  sc                   383 non-null    float64
13  sod                  313 non-null    float64
14  pot                  312 non-null    float64
15  hemo                 348 non-null    float64
16  pcv                  330 non-null    object
17  wc                   295 non-null    object
18  rc                   270 non-null    object
19  htn                  398 non-null    object
20  dm                   398 non-null    object
21  cad                  398 non-null    object
22  appet               399 non-null    object
23  pe                   399 non-null    object
24  ane                  399 non-null    object
25  classification       400 non-null    object
dtypes: float64(11), int64(1), object(14)
memory usage: 81.4+ KB
```

In [6]:

```
kidneyDisease.describe()
```

Out[6]:

	id	age	bp	sg	al	su	bgr	
<b>count</b>	400.000000	391.000000	388.000000	353.000000	354.000000	351.000000	356.000000	381
<b>mean</b>	199.500000	51.483376	76.469072	1.017408	1.016949	0.450142	148.036517	57
<b>std</b>	115.614301	17.169714	13.683637	0.005717	1.352679	1.099191	79.281714	50
<b>min</b>	0.000000	2.000000	50.000000	1.005000	0.000000	0.000000	22.000000	1
<b>25%</b>	99.750000	42.000000	70.000000	1.010000	0.000000	0.000000	99.000000	27
<b>50%</b>	199.500000	55.000000	80.000000	1.020000	0.000000	0.000000	121.000000	42
<b>75%</b>	299.250000	64.500000	80.000000	1.020000	2.000000	0.000000	163.000000	66
<b>max</b>	399.000000	90.000000	180.000000	1.025000	5.000000	5.000000	490.000000	391

# Data pre-processing

In [7]:

```
cols_names={"bp":"blood_pressure",
            "sg":"specific_gravity",
            "al":"albumin",
            "su":"sugar",
            "rbc":"red_blood_cells",
            "pc":"pus_cell",
            "pcc":"pus_cell_clumps",
            "ba":"bacteria",
            "bgr":"blood_glucose_random",
            "bu":"blood_urea",
            "sc":"serum_creatinine",
            "sod":"sodium",
            "pot":"potassium",
            "hemo":"haemoglobin",
            "pcv":"packed_cell_volume",
            "wc":"white_blood_cell_count",
            "rc":"red_blood_cell_count",
            "htn":"hypertension",
            "dm":"diabetes_mellitus",
            "cad":"coronary_artery_disease",
            "appet":"appetite",
            "pe":"pedal_edema",
            "ane":"anemia"}

kidneyDisease.rename(columns=cols_names, inplace=True)
```

In [8]:

```
# Dropping the id column and customization

kidneyDisease['red_blood_cell_count'] = pd.to_numeric(kidneyDisease['red_blood_cell_count'],
                                                       errors='coerce')
kidneyDisease['packed_cell_volume'] = pd.to_numeric(kidneyDisease['packed_cell_volume'],
                                                      errors='coerce')
kidneyDisease['white_blood_cell_count'] = pd.to_numeric(kidneyDisease['white_blood_cell_count'],
                                                         errors='coerce')

kidneyDisease.drop(["id"],axis=1,inplace=True)
```

In [9]:

```
kidneyDisease.isnull().sum().sort_values(ascending=False)
```

Out[9]:

red_blood_cells	152
red_blood_cell_count	131
white_blood_cell_count	106
potassium	88
sodium	87
packed_cell_volume	71
pus_cell	65
haemoglobin	52
sugar	49
specific_gravity	47
albumin	46
blood_glucose_random	44
blood_urea	19
serum_creatinine	17
blood_pressure	12
age	9
bacteria	4
pus_cell_clumps	4
hypertension	2
diabetes_mellitus	2
coronary_artery_disease	2
anemia	1
appetite	1
pedal_edema	1
classification	0
dtype: int64	

In [10]:

```
# Numerical & Categorical features:
```

In [11]:

```
numerical_features = []
categorical_features = []

for i in kidneyDisease.columns:
    if kidneyDisease[i].nunique()>7:
        numerical_features.append(i)
    else:
        categorical_features.append(i)

# Numerical
print(numerical_features)
print("\n")
# Categorical
print(categorical_features)

['age', 'blood_pressure', 'blood_glucose_random', 'blood_urea', 'serum_creatinine', 'sodium', 'potassium', 'haemoglobin', 'packed_cell_volume', 'white_blood_cell_count', 'red_blood_cell_count']

['specific_gravity', 'albumin', 'sugar', 'red_blood_cells', 'pus_cell', 'pus_cell_clumps', 'bacteria', 'hypertension', 'diabetes_mellitus', 'coronary_artery_disease', 'appetite', 'pedal_edema', 'anemia', 'classification']
```

In [12]:

```
# Replace incorrect values

kidneyDisease['diabetes_mellitus'] = kidneyDisease['diabetes_mellitus'].replace(to_replace='diabetes_mellitus', value='diabetes_mellitus')

kidneyDisease['coronary_artery_disease'] = kidneyDisease['coronary_artery_disease'].replace(to_replace='coronary_artery_disease', value='coronary_artery_disease')

kidneyDisease['classification'] = kidneyDisease['classification'].replace(to_replace='classification', value='classification')
```

Actual work

In [13]:

```
# Duplicating df
df = kidneyDisease

# Filter the Numerical data
df1 = df.filter(numerical_features, axis=1)
```

In [14]:

```
# Case 1 df with NOT NULL Original data
df_notnull = df.dropna().filter(numerical_features, axis=1)
df_notnull
```

Out[14]:

	age	blood_pressure	blood_glucose_random	blood_urea	serum_creatinine	sodium	potass
<b>3</b>	48.0	70.0	117.0	56.0	3.8	111.0	
<b>9</b>	53.0	90.0	70.0	107.0	7.2	114.0	
<b>11</b>	63.0	70.0	380.0	60.0	2.7	131.0	
<b>14</b>	68.0	80.0	157.0	90.0	4.1	130.0	
<b>20</b>	61.0	80.0	173.0	148.0	3.9	135.0	
...	...	...	...	...	...	...	
<b>395</b>	55.0	80.0	140.0	49.0	0.5	150.0	
<b>396</b>	42.0	70.0	75.0	31.0	1.2	141.0	
<b>397</b>	12.0	80.0	100.0	26.0	0.6	137.0	
<b>398</b>	17.0	60.0	114.0	50.0	1.0	135.0	
<b>399</b>	58.0	80.0	131.0	18.0	1.1	141.0	

158 rows × 11 columns

In [15]:

```
# Case 2 df with Imputed mean data
```

```
from sklearn.impute import SimpleImputer
miss_mean_imputer = SimpleImputer(missing_values=np.nan, strategy='mean')
miss_mean_imputer = miss_mean_imputer.fit(df1)
df_imputed = miss_mean_imputer.transform(df1.values)
df_imputed = pd.DataFrame(df_imputed, columns = numerical_features)
df_imputed
```

Out[15]:

	age	blood_pressure	blood_glucose_random	blood_urea	serum_creatinine	sodium	poi
0	48.0	80.0	121.000000	36.0	1.2	137.528754	4
1	7.0	50.0	148.036517	18.0	0.8	137.528754	4
2	62.0	80.0	423.000000	53.0	1.8	137.528754	4
3	48.0	70.0	117.000000	56.0	3.8	111.000000	2
4	51.0	80.0	106.000000	26.0	1.4	137.528754	4
...	...	...	...	...	...	...	...
395	55.0	80.0	140.000000	49.0	0.5	150.000000	4
396	42.0	70.0	75.000000	31.0	1.2	141.000000	3
397	12.0	80.0	100.000000	26.0	0.6	137.000000	4
398	17.0	60.0	114.000000	50.0	1.0	135.000000	4
399	58.0	80.0	131.000000	18.0	1.1	141.000000	3

400 rows × 11 columns


In [16]:

```
#Jointplot
```


```
sns.jointplot(x='age',y='blood_glucose_random',data=df_notnull, color='brown')  
df_notnull = df_notnull[df_notnull['blood_glucose_random'].between(0, 300)]  
sns.lmplot(x='age',y='blood_glucose_random',data=df_notnull)
```

Out[16]:

<seaborn.axisgrid.FacetGrid at 0x7feae11b3550>








In [17]:



```
#Jointplot
sns.jointplot(x='age',y='blood_urea',data=df_notnull,color='brown')
df_notnull = df_notnull[df_notnull['blood_urea'].between(0, 120)]
sns.lmplot(x='age',y='blood_urea',data=df_notnull)

sns.jointplot(x='age',y='blood_urea',data=df_imputed,color='brown')
df_imputed = df_imputed[df_imputed['blood_urea'].between(0, 120)]
sns.lmplot(x='age',y='blood_urea',data=df_imputed)
# blood_glucose_random blood_urea serum_creatinine sodium
```

Out[17]:

<seaborn.axisgrid.FacetGrid at 0x7feae1979130>





In [18]:


```
#Jointplot
sns.jointplot(x='age',y='serum_creatinine',data=df_notnull,color='brown')
df_notnull = df_notnull[df_notnull['serum_creatinine'].between(0, 5)]
sns.lmplot(x='age',y='serum_creatinine',data=df_notnull)



sns.jointplot(x='age',y='serum_creatinine',data=df_imputed,color='brown')
df_imputed = df_imputed[df_imputed['serum_creatinine'].between(0, 5)]
sns.lmplot(x='age',y='serum_creatinine',data=df_imputed)


# blood_glucose_random blood_urea serum_creatinine sodium
```

Out[18]:

<seaborn.axisgrid.FacetGrid at 0x7feae17497c0>







In [19]:


```
#Jointplot
```


```
sns.jointplot(x='age',y='sodium',data=df_imputed,color='brown')  
df_imputed = df_imputed[df_imputed['sodium'].between(80, 180)]  
sns.lmplot(x='age',y='sodium',data=df_imputed)
```

```
# blood_glucose_random blood_urea serum_creatinine sodium
```

Out[19]:

<seaborn.axisgrid.FacetGrid at 0x7feae12c9940>







In [20]:



```
#Jointplot
```

```
sns.jointplot(x='age',y='potassium',data=df_imputed,color='brown')  
sns.lmplot(x='age',y='potassium',data=df_imputed)
```

```
# blood_glucose_random blood_urea serum_creatinine sodium potassium
```

Out[20]:

<seaborn.axisgrid.FacetGrid at 0x7feae2ab96d0>



In [21]:



```
#Jointplot
sns.jointplot(x='age',y='haemoglobin',data=df_notnull,color='brown')
sns.lmplot(x='age',y='haemoglobin',data=df_notnull)



sns.jointplot(x='age',y='haemoglobin',data=df_imputed,color='brown')
sns.lmplot(x='age',y='haemoglobin',data=df_imputed)

# blood_glucose_random blood_urea serum_creatinine sodium potassium haemoglobin
```

Out[21]:

<seaborn.axisgrid.FacetGrid at 0x7feae1542220>






In [22]:

```
sns.pairplot(df1,  
             x_vars = ['blood_urea', 'blood_pressure',  
                       'haemoglobin', 'sodium'],  
             y_vars = 'age',  
             kind = "reg" )
```

Out[22]:

<seaborn.axisgrid.PairGrid at 0x7feae163d250>

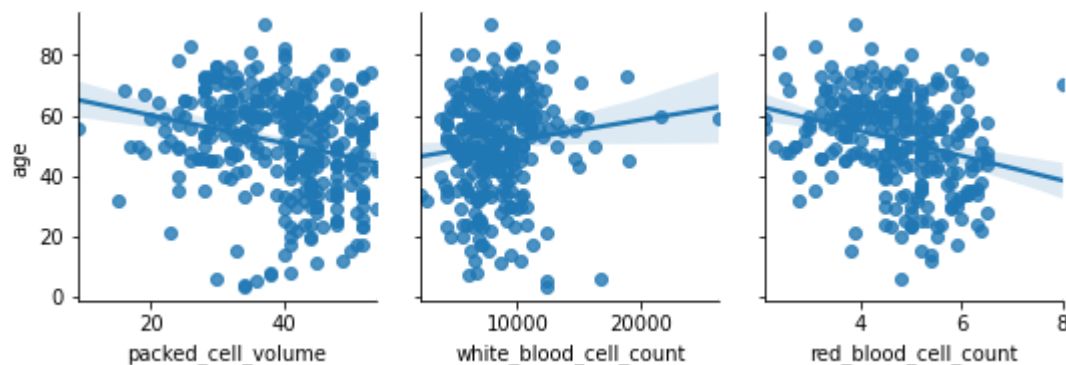


In [23]:

```
sns.pairplot(df1,  
             x_vars = ['packed_cell_volume',  
                       'white_blood_cell_count', 'red_blood_cell_count'],  
             y_vars = 'age',  
             kind = "reg" )
```

Out[23]:

<seaborn.axisgrid.PairGrid at 0x7feae0e28700>



In [24]:

```
sns.pairplot(df_notnull)
```

In [25]:

```
sns.pairplot(df_imputed)
```

## Modelling

### Simple Linear Regression

In [94]:

```
# blood_glucose_random blood_urea serum_creatinine sodium potassium haemoglobin packed_cell_volume
X = df_notnull[["packed_cell_volume"]]
Y = df_notnull[["haemoglobin"]]
```

In [95]:

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, Y) #, test_size=0.3, random_state=42

from sklearn.linear_model import LinearRegression
linreg = LinearRegression()
linreg.fit(X_train, y_train)
```

Out[95]:

LinearRegression()

In [96]:

```
#Training Accuracy
linreg.score(X_train,y_train)

#Prediction
prediction=linreg.predict(X_test)

#Testing Accuracy
linreg.score(X_test,y_test)
```

Out[96]:


0.5701234680464349

In [97]:

```
plt.scatter(y_test,prediction)
```

Out[97]:

<matplotlib.collections.PathCollection at 0x7feae2aed760>



In [98]:

```
from sklearn import metrics
print("MAE: ",metrics.mean_absolute_error(y_test,prediction))
print("MSE: ",metrics.mean_squared_error(y_test,prediction))
print("RMSE: ",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
MAE:  1.31760321006809
MSE:  2.454089260358441
RMSE:  1.5665533059422017
```

In [99]:

```
coef=pd.DataFrame()
coef['Features'] = X.columns.values
coef['Coefficients'] = linreg.coef_
coef
```

Out[99]:

	Features	Coefficients
0	packed_cell_volume	0.213976

In [100]:

```
#Testing Accuracy
linreg.score(X_test,y_test)
```

Out[100]:

```
0.5701234680464349
```

## Multiple Linear Regression

In [101]:

```
#Splitting into Training and Testing Data
X=df_notnull.drop(['age'],axis=1)
y=df_notnull[['red_blood_cell_count','potassium']]
X
```

Out[101]:

	blood_pressure	blood_glucose_random	blood_urea	serum_creatinine	sodium	potassium
3	70.0	117.0	56.0	3.8	111.0	2.5
14	80.0	157.0	90.0	4.1	130.0	6.4
27	70.0	264.0	87.0	2.7	130.0	4.0
48	70.0	70.0	32.0	0.9	125.0	4.0
71	60.0	163.0	92.0	3.3	141.0	4.0
...	...	...	...	...	...	...
395	80.0	140.0	49.0	0.5	150.0	4.9
396	70.0	75.0	31.0	1.2	141.0	3.5
397	80.0	100.0	26.0	0.6	137.0	4.4
398	60.0	114.0	50.0	1.0	135.0	4.9
399	80.0	131.0	18.0	1.1	141.0	3.5

133 rows × 10 columns

In [102]:

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.3, random_stat
```

In [103]:

```
from sklearn.linear_model import LinearRegression

multi_lr=LinearRegression()
multi_lr.fit(X_train,y_train)

#Training Accuracy
multi_lr.score(X_train,y_train)
```

Out[103]:

1.0

In [104]:


```
#Prediction
prediction=multi_lr.predict(X_test)
```

In [105]:

```
plt.scatter(y_test,prediction)
```

Out[105]:

<matplotlib.collections.PathCollection at 0x7feae160adf0>



In [106]:

```
from sklearn import metrics
print("MAE: ",metrics.mean_absolute_error(y_test,prediction))
print("MSE: ",metrics.mean_squared_error(y_test,prediction))
print("RMSE: ",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
MAE:  2.7977620220553944e-15
MSE:  1.6487192919119146e-29
RMSE:  4.060442453615018e-15
```



In [107]:

```
#Splitting into Training and Testing Data
X=df_notnull.drop(['red_blood_cell_count','potassium','packed_cell_volume','blood_ur
y=df_notnull[['age']]
X
```

Out[107]:

	age	blood_pressure	blood_glucose_random	serum_creatinine	sodium	haemoglobin	white
3	48.0	70.0	117.0	3.8	111.0	11.2	
14	68.0	80.0	157.0	4.1	130.0	5.6	
27	69.0	70.0	264.0	2.7	130.0	12.5	
48	73.0	70.0	70.0	0.9	125.0	10.0	
71	46.0	60.0	163.0	3.3	141.0	9.8	
...	...	...	...	...	...	...	
395	55.0	80.0	140.0	0.5	150.0	15.7	
396	42.0	70.0	75.0	1.2	141.0	16.5	
397	12.0	80.0	100.0	0.6	137.0	15.8	
398	17.0	60.0	114.0	1.0	135.0	14.2	
399	58.0	80.0	131.0	1.1	141.0	15.8	

133 rows × 7 columns

In [108]:

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.3, random_stat
```

In [109]:

```
from sklearn.linear_model import LinearRegression

multi_lr=LinearRegression()
multi_lr.fit(X_train,y_train)

#Training Accuracy
multi_lr.score(X_train,y_train)
```

Out[109]:

1.0

In [110]:


```
#Prediction
prediction=multi_lr.predict(X_test)
```

In [111]:

```
plt.scatter(y_test,prediction)
```

Out[111]:

<matplotlib.collections.PathCollection at 0x7feae2460f10>



In [112]:

```
from sklearn import metrics
print("MAE: ",metrics.mean_absolute_error(y_test,prediction))
print("MSE: ",metrics.mean_squared_error(y_test,prediction))
print("RMSE: ",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
MAE: 2.930988785010413e-15
MSE: 2.6979042958558606e-29
RMSE: 5.194135438988726e-15
```

## Inference

Linear regression is one of the most common techniques of regression analysis. It is also called a simple linear regression. It establishes the relationship between two variables using a straight line. Linear regression attempts to draw a line that comes closest to the data by finding the slope and intercept that define the line and minimize regression errors.

Multiple regression is a broader class of regressions that encompasses linear and nonlinear regressions with multiple explanatory variables.

It is rare that a dependent variable is explained by only one variable. In this case, an analyst uses multiple regression, which attempts to explain dependent variable using more than one independent variable. Multiple regressions can be linear and nonlinear.

In [ ]:

