

# MANOJ KUMAR - 2048015

## Lab 10: Naive Bayes (NB) Classification

Kidney diseases are disorders that affect the functions of the kidney. During the late stages, kidney diseases can cause kidney failure to prevent chronic kidney disease-CKD by utilising machine learning techniques to diagnose kidney disease at an early stage. We describe the most prominent supervised machine learning algorithms (SML), their characteristics, Generalisation capacity of each method, Time complexity, Hyper-parameter tuning, and Advantages and disadvantages of each technique comparatively.

The Kidney Disease dataset obtained from UCI was used to determine and test its highest percentage of accuracy and benchmark.

### *Importing basic libraries*

```
In [1]: import pandas as pd
import numpy as np
import time
import seaborn as sns
import matplotlib.pyplot as plt
import matplotlib as mpl
```

### *Reading the dataset*

```
In [2]: ckd_df = pd.read_csv('kidney_disease.csv')

#Check the shape
print(ckd_df.shape)

(400, 26)
```

```
In [3]: #check the columns
ckd_df.columns
```

```
Out[3]: Index(['id', 'age', 'bp', 'sg', 'al', 'su', 'rbc', 'pc', 'pcc', 'ba', 'bgr',
              'bu', 'sc', 'sod', 'pot', 'hemo', 'pcv', 'wc', 'rc', 'htn', 'dm', 'cad',
              'appet', 'pe', 'ane', 'classification'],
              dtype='object')
```

### *Rename the columns to have meaningful names*

```
In [4]: col_dict={"bp": "blood_pressure",
                  "sg": "specific_gravity",
                  "al": "albumin",
                  "su": "sugar",
                  "rbc": "red_blood_cells",
                  "pc": "pus_cell",
                  "pcc": "pus_cell_clumps",
                  "ba": "bacteria",
                  "bgr": "blood_glucose_random",
                  "bu": "blood_urea",
                  "sc": "serum_creatinine",
                  "sod": "sodium",
                  "pot": "potassium",
                  "hemo": "hemoglobin",
                  "pcv": "packed_cell_volume",
                  "wc": "white_blood_cell_count",
                  "rc": "red_blood_cell_count",
                  "htn": "hypertension",
                  "dm": "diabetes_mellitus",
                  "cad": "coronary_artery_disease",
                  "appet": "appetite",
                  "pe": "pedal_edema",
                  "ane": "anemia"}

ckd_df.rename(columns=col_dict, inplace=True)

#Check the column names again
ckd_df.columns
```

### Observing the data

```
In [5]: ckd_df.head(11).T
```

[illegible]

```
In [6]: for i in ckd_df.drop("id",axis=1).columns:
        print('unique values in "{}":\n'.format(i),ckd_df[i].unique())

unique values in "age":
[48.  7. 62. 51. 60. 68. 24. 52. 53. 50. 63. 40. 47. 61. 21. 42. 75. 69.
 nan 73. 70. 65. 76. 72. 82. 46. 45. 35. 54. 11. 59. 67. 15. 55. 44. 26.
 64. 56.  5. 74. 38. 58. 71. 34. 17. 12. 43. 41. 57.  8. 39. 66. 81. 14.
 27. 83. 30.  4.  3.  6. 32. 80. 49. 90. 78. 19.  2. 33. 36. 37. 23. 25.
 20. 29. 28. 22. 79.]
unique values in "blood_pressure":
[ 80.  50.  70.  90.  nan 100.  60. 110. 140. 180. 120.]
unique values in "specific_gravity":
[1.02  1.01  1.005 1.015  nan 1.025]
unique values in "albumin":
[ 1.  4.  2.  3.  0. nan  5.]
unique values in "sugar":
[ 0.  3.  4.  1. nan  2.  5.]
unique values in "red_blood_cells":
[nan 'normal' 'abnormal']
unique values in "pus_cell":
['normal' 'abnormal' nan]
unique values in "pus_cell_clumps":
[ 0.  1.  2.  3.  4.  5.]
```

```
In [7]: #Replace incorrect values
ckd_df['diabetes_mellitus'] =ckd_df['diabetes_mellitus'].replace(to_replace={'\tno':'no','\tyes':'yes',' yes':''})
ckd_df['coronary_artery_disease'] = ckd_df['coronary_artery_disease'].replace(to_replace='\tno',value='no')
ckd_df['white_blood_cell_count'] = ckd_df['white_blood_cell_count'].replace(to_replace='\t8400',value='8400')
ckd_df["classification"]=ckd_df["classification"].replace("ckd\t", "ckd")

for i in range(ckd_df.shape[0]):
    if ckd_df.iloc[i,16]=='\t?':
        ckd_df.iloc[i,16]=np.nan
    if ckd_df.iloc[i,16]=='\t43':
        ckd_df.iloc[i,16]='43'
    if ckd_df.iloc[i,17]=='\t?':
        ckd_df.iloc[i,17]=np.nan
    if ckd_df.iloc[i,17]=='\t6200':
        ckd_df.iloc[i,17]='6200'
    if ckd_df.iloc[i,18]=='\t?':
        ckd_df.iloc[i,18]=np.nan
    if ckd_df.iloc[i,25]=='ckd':
        ckd_df.iloc[i,25]='1'
    if ckd_df.iloc[i,25]=='notckd':
        ckd_df.iloc[i,25]='0'

for i in ckd_df.drop("id",axis=1).columns:
    print('unique values in "{}":\n'.format(i),ckd_df[i].unique())
```

```
unique values in "age":
[48.  7. 62. 51. 60. 68. 24. 52. 53. 50. 63. 40. 47. 61. 21. 42. 75. 69.
 nan 73. 70. 65. 76. 72. 82. 46. 45. 35. 54. 11. 59. 67. 15. 55. 44. 26.
 64. 56.  5. 74. 38. 58. 71. 34. 17. 12. 43. 41. 57.  8. 39. 66. 81. 14.
 27. 83. 30.  4.  3.  6. 32. 80. 49. 90. 78. 19.  2. 33. 36. 37. 23. 25.
 20. 29. 28. 22. 79.]
unique values in "blood_pressure":
[ 80.  50.  70.  90.  nan 100.  60. 110. 140. 180. 120.]
unique values in "specific_gravity":
[1.02  1.01  1.005 1.015  nan 1.025]
unique values in "albumin":
[ 1.  4.  2.  3.  0. nan  5.]
unique values in "sugar":
[ 0.  3.  4.  1. nan  2.  5.]
unique values in "red_blood_cells":
[nan 'normal' 'abnormal']
unique values in "pus_cell":
['normal' 'abnormal' nan]
unique values in "pus_cell_clumps":
['notpresent' 'present' nan]
unique values in "bacteria":
['notpresent' 'present' nan]
unique values in "blood_glucose_random":
[121.  nan 423. 117. 106.  74. 100. 410. 138.  70. 490. 380. 208.  98.
 157.  76.  99. 114. 263. 173.  95. 108. 156. 264. 123.  93. 107. 159.
 140. 171. 270.  92. 137. 204.  79. 207. 124. 144.  91. 162. 246. 253.
 141. 182.  86. 150. 146. 425. 112. 250. 360. 163. 129. 133. 102. 158.
 165. 132. 104. 127. 415. 169. 251. 109. 280. 210. 219. 295.  94. 172.
 101. 298. 153.  88. 226. 143. 115.  89. 297. 233. 294. 323. 125.  90.
 308. 118. 224. 128. 122. 214. 213. 268. 256.  84. 105. 288. 139.  78.
 273. 242. 424. 303. 148. 160. 192. 307. 220. 447. 309.  22. 111. 261.
 215. 234. 131. 352.  80. 239. 110. 130. 184. 252. 113. 230. 341. 255.
 103. 238. 248. 120. 241. 269. 201. 203. 463. 176.  82. 119.  97.  96.
  81. 116. 134.  85.  83.  87.  75.]
unique values in "blood_urea":
[ 36.  18.  53.  56.  26.  25.  54.  31.  60. 107.  55.  72.
  86.  90. 162.  46.  87.  27. 148. 180. 163.  nan  50.  75.
  45.  28. 155.  33.  39. 153.  29.  65. 103.  70.  80.  20.
 202.  77.  89.  24.  17.  32. 114.  66.  38. 164. 142.  96.
 391.  15. 111.  73.  19.  92.  35.  16. 139.  48.  85.  98.
 186.  37.  47.  52.  82.  51. 106.  22. 217.  88. 118.  50.1
  71.  34.  40.  21. 219.  30. 125. 166.  49. 208. 176.  68.
 145. 165. 322.  23. 235. 132.  76.  42.  44.  41. 113.  1.5
 146.  58. 133. 137.  67. 115. 223.  98.6 158.  94.  74. 150.
  61.  57.  95. 191.  93. 241.  64.  79. 215. 309.  10. ]
unique values in "serum_creatinine":
[ 1.2  0.8  1.8  3.8  1.4  1.1  24.  1.9  7.2  4.  2.7  2.1
  4.6  4.1  9.6  2.2  5.2  1.3  1.6  3.9 76.  7.7  nan  2.4
  7.3  1.5  2.5  2.  3.4  0.7  1. 10.8  6.3  5.9  0.9  3.
  3.25 9.7  6.4  3.2 32.  0.6  6.1  3.3  6.7  8.5  2.8 15.
  2.9  1.7  3.6  5.6  6.5  4.4 10.2 11.5  0.5 12.2  5.3  9.2
 13.8 16.9  6.  7.1 18.  2.3 13.  48.1 14.2 16.4  2.6  7.5
  4.3 18.1 11.8  9.3  6.8 13.5 12.8 11.9 12.  13.4 15.2 13.3
 0.4 ]
unique values in "sodium":
[ nan 111. 142. 104. 114. 131. 138. 135. 130. 141. 139.  4.5
 136. 129. 140. 132. 133. 134. 125. 163. 137. 128. 143. 127.
 146. 126. 122. 147. 124. 115. 145. 113. 120. 150. 144. ]
unique values in "potassium":
[ nan 2.5 3.2  4.  3.7 4.2  5.8 3.4  6.4 4.9 4.1 4.3  5.2 3.8
  4.6 3.9 4.7  5.9 4.8 4.4  6.6 39.  5.5  5.  3.5 3.6  7.6 2.9
  4.5 5.7  5.4  5.3 47.  6.3  5.1  5.6  3.  2.8 2.7  6.5 3.3]
unique values in "hemoglobin":
```

```

[15.4 11.3  9.6 11.2 11.6 12.2 12.4 10.8  9.5  9.4  9.7  9.8  5.6  7.6
12.6 12.1 12.7 10.3  7.7 10.9  nan 11.1  9.9 12.5 12.9 10.1 12.  13.
 7.9  9.3 15.  10.   8.6 13.6 10.2 10.5  6.6 11.   7.5 15.6 15.2  4.8
 9.1  8.1 11.9 13.5  8.3  7.1 16.1 10.4  9.2  6.2 13.9 14.1  6.  11.8
11.7 11.4 14.   8.2 13.2  6.1  8.  12.3  8.4 14.3  9.   8.7 10.6 13.1
10.7  5.5  5.8  6.8  8.8  8.5 13.8 11.5  7.3 13.7 12.8 13.4  6.3  3.1
17.  15.9 14.5 15.5 16.2 14.4 14.2 16.3 14.8 16.5 15.7 13.3 14.6 16.4
16.9 16.  14.7 16.6 14.9 16.7 16.8 15.8 15.1 17.1 17.2 15.3 17.3 17.4
17.7 17.8 17.5 17.6]
unique values in "packed_cell_volume":
['44' '38' '31' '32' '35' '39' '36' '33' '29' '28' nan '16' '24' '37' '30'
'34' '40' '45' '27' '48' '52' '14' '22' '18' '42' '17' '46' '23' '19'
'25' '41' '26' '15' '21' '43' '20' '47' '9' '49' '50' '53' '51' '54']
unique values in "white_blood_cell_count":
['7800' '6000' '7500' '6700' '7300' nan '6900' '9600' '12100' '4500'
'12200' '11000' '3800' '11400' '5300' '9200' '6200' '8300' '8400' '10300'
'9800' '9100' '7900' '6400' '8600' '18900' '21600' '4300' '8500' '11300'
'7200' '7700' '14600' '6300' '7100' '11800' '9400' '5500' '5800' '13200'
'12500' '5600' '7000' '11900' '10400' '10700' '12700' '6800' '6500'
'13600' '10200' '9000' '14900' '8200' '15200' '5000' '16300' '12400'
'10500' '4200' '4700' '10900' '8100' '9500' '2200' '12800' '11200'
'19100' '12300' '16700' '2600' '26400' '8800' '7400' '4900' '8000'
'12000' '15700' '4100' '5700' '11500' '5400' '10800' '9900' '5200' '5900'
'9300' '9700' '5100' '6600']
unique values in "red_blood_cell_count":
['5.2' nan '3.9' '4.6' '4.4' '5' '4.0' '3.7' '3.8' '3.4' '2.6' '2.8' '4.3'
'3.2' '3.6' '4' '4.1' '4.9' '2.5' '4.2' '4.5' '3.1' '4.7' '3.5' '6.0'
'5.0' '2.1' '5.6' '2.3' '2.9' '2.7' '8.0' '3.3' '3.0' '3' '2.4' '4.8'
'5.4' '6.1' '6.2' '6.3' '5.1' '5.8' '5.5' '5.3' '6.4' '5.7' '5.9' '6.5']
unique values in "hypertension":
['yes' 'no' nan]
unique values in "diabetes_mellitus":
['yes' 'no' nan]
unique values in "coronary_artery_disease":
['no' 'yes' nan]
unique values in "appetite":
['good' 'poor' nan]
unique values in "pedal_edema":
['no' 'yes' nan]
unique values in "anemia":
['no' 'yes' nan]
unique values in "classification":
['1' '0']

```

```

In [8]: # Observing the summarized information of data
ckd_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 26 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     400 non-null    int64
1   age                                   391 non-null    float64
2   blood_pressure                       388 non-null    float64
3   specific_gravity                    353 non-null    float64
4   albumin                             354 non-null    float64
5   sugar                               351 non-null    float64
6   red_blood_cells                     248 non-null    object
7   pus_cell                            335 non-null    object
8   pus_cell_clumps                     396 non-null    object
9   bacteria                            396 non-null    object
10  blood_glucose_random                 356 non-null    float64
11  blood_urea                           381 non-null    float64
12  serum_creatinine                    383 non-null    float64
13  sodium                              313 non-null    float64
14  potassium                            312 non-null    float64
15  hemoglobin                           348 non-null    float64
16  packed_cell_volume                  329 non-null    object
17  white_blood_cell_count              294 non-null    object
18  red_blood_cell_count                269 non-null    object
19  hypertension                         398 non-null    object
20  diabetes_mellitus                   398 non-null    object
21  coronary_artery_disease             398 non-null    object
22  appetite                            399 non-null    object
23  pedal_edema                         399 non-null    object
24  anemia                              399 non-null    object
25  classification                       400 non-null    object
dtypes: float64(11), int64(1), object(14)
memory usage: 81.4+ KB

```

```
In [9]: ckd_df.iloc[:, -1]=ckd_df.iloc[:, -1].astype('int64')
ckd_df.head(11).T
```

Out[9]:

	0	1	2	3	4	5	6	7	8	9	10
id	0	1	2	3	4	5	6	7	8	9	10
age	48	7	62	48	51	60	68	24	52	53	50
blood_pressure	80	50	80	70	80	90	70	NaN	100	90	60
specific_gravity	1.02	1.02	1.01	1.005	1.01	1.015	1.01	1.015	1.015	1.02	1.01
albumin	1	4	2	4	2	3	0	2	3	2	2
sugar	0	0	3	0	0	0	0	4	0	0	4
red_blood_cells	NaN	NaN	normal	normal	normal	NaN	NaN	normal	normal	abnormal	NaN
pus_cell	normal	normal	normal	abnormal	normal	NaN	normal	abnormal	abnormal	abnormal	abnormal
pus_cell_clumps	notpresent	notpresent	notpresent	present	notpresent	notpresent	notpresent	notpresent	present	present	present
bacteria	notpresent	notpresent	notpresent	notpresent	notpresent	notpresent	notpresent	notpresent	notpresent	notpresent	notpresent
blood_glucose_random	121	NaN	423	117	106	74	100	410	138	70	490
blood_urea	36	18	53	56	26	25	54	31	60	107	55
serum_creatinine	1.2	0.8	1.8	3.8	1.4	1.1	24	1.1	1.9	7.2	4
sodium	NaN	NaN	NaN	111	NaN	142	104	NaN	NaN	114	NaN
potassium	NaN	NaN	NaN	2.5	NaN	3.2	4	NaN	NaN	3.7	NaN
hemoglobin	15.4	11.3	9.6	11.2	11.6	12.2	12.4	12.4	10.8	9.5	9.4
packed_cell_volume	44	38	31	32	35	39	36	44	33	29	28
white_blood_cell_count	7800	6000	7500	6700	7300	7800	NaN	6900	9600	12100	NaN
red_blood_cell_count	5.2	NaN	NaN	3.9	4.6	4.4	NaN	5	4.0	3.7	NaN
hypertension	yes	no	no	yes	no	yes	no	no	yes	yes	yes
diabetes_mellitus	yes	no	yes	no	no	yes	no	yes	yes	yes	yes
coronary_artery_disease	no	no	no	no	no	no	no	no	no	no	no
appetite	good	good	poor	poor	good	good	good	good	good	poor	good
pedal_edema	no	no	no	yes	no	yes	no	yes	no	no	no
anemia	no	no	yes	yes	no	no	no	no	yes	yes	yes
classification	1	1	1	1	1	1	1	1	1	1	1

```
In [10]: print(ckd_df['packed_cell_volume'].unique())
print(ckd_df['white_blood_cell_count'].unique())
print(ckd_df['red_blood_cell_count'].unique())
```

['44' '38' '31' '32' '35' '39' '36' '33' '29' '28' nan '16' '24' '37' '30' '34' '40' '45' '27' '48' '52' '14' '22' '18' '42' '17' '46' '23' '19' '25' '41' '26' '15' '21' '43' '20' '47' '9' '49' '50' '53' '51' '54']

['7800' '6000' '7500' '6700' '7300' nan '6900' '9600' '12100' '4500' '12200' '11000' '3800' '11400' '5300' '9200' '6200' '8300' '8400' '10300' '9800' '9100' '7900' '6400' '8600' '18900' '21600' '4300' '8500' '11300' '7200' '7700' '14600' '6300' '7100' '11800' '9400' '5500' '5800' '13200' '12500' '5600' '7000' '11900' '10400' '10700' '12700' '6800' '6500' '13600' '10200' '9000' '14900' '8200' '15200' '5000' '16300' '12400' '10500' '4200' '4700' '10900' '8100' '9500' '2200' '12800' '11200' '19100' '12300' '16700' '2600' '26400' '8800' '7400' '4900' '8000' '12000' '15700' '4100' '5700' '11500' '5400' '10800' '9900' '5200' '5900' '9300' '9700' '5100' '6600']

['5.2' nan '3.9' '4.6' '4.4' '5' '4.0' '3.7' '3.8' '3.4' '2.6' '2.8' '4.3' '3.2' '3.6' '4' '4.1' '4.9' '2.5' '4.2' '4.5' '3.1' '4.7' '3.5' '6.0' '5.0' '2.1' '5.6' '2.3' '2.9' '2.7' '8.0' '3.3' '3.0' '3' '2.4' '4.8' '5.4' '6.1' '6.2' '6.3' '5.1' '5.8' '5.5' '5.3' '6.4' '5.7' '5.9' '6.5']

```
In [11]: mistyped=['packed_cell_volume','white_blood_cell_count','red_blood_cell_count']
        for col in mistyped:
            ckd_df[col]=ckd_df[col].astype('float')

        numeric=[]
        for i in ckd_df.columns:
            if ckd_df[i].dtype=='float64':
                numeric.append(i)

        numeric
```

```
Out[11]: ['age',
          'blood_pressure',
          'specific_gravity',
          'albumin',
          'sugar',
          'blood_glucose_random',
          'blood_urea',
          'serum_creatinine',
          'sodium',
          'potassium',
          'hemoglobin',
          'packed_cell_volume',
          'white_blood_cell_count',
          'red_blood_cell_count']
```

```
In [12]: ckd_df.drop('id',axis=1,inplace=True)

        categoricals=[]

        for col in ckd_df.columns:
            if not col in numeric:
                categoricals.append(col)
        categoricals.remove('classification')

        categoricals
```

```
Out[12]: ['red_blood_cells',
          'pus_cell',
          'pus_cell_clumps',
          'bacteria',
          'hypertension',
          'diabetes_mellitus',
          'coronary_artery_disease',
          'appetite',
          'pedal_edema',
          'anemia']
```

```
In [13]: import warnings
        warnings.simplefilter('ignore')

        import matplotlib.style as style
        style.use('fivethirtyeight')
```

**Checking distribution of the numerical features**

```
In [14]: fig, axes = plt.subplots(nrows=7, ncols=2, figsize=(15,30))
fig.subplots_adjust(hspace=0.5)
fig.suptitle('Distributions of numerical Features')

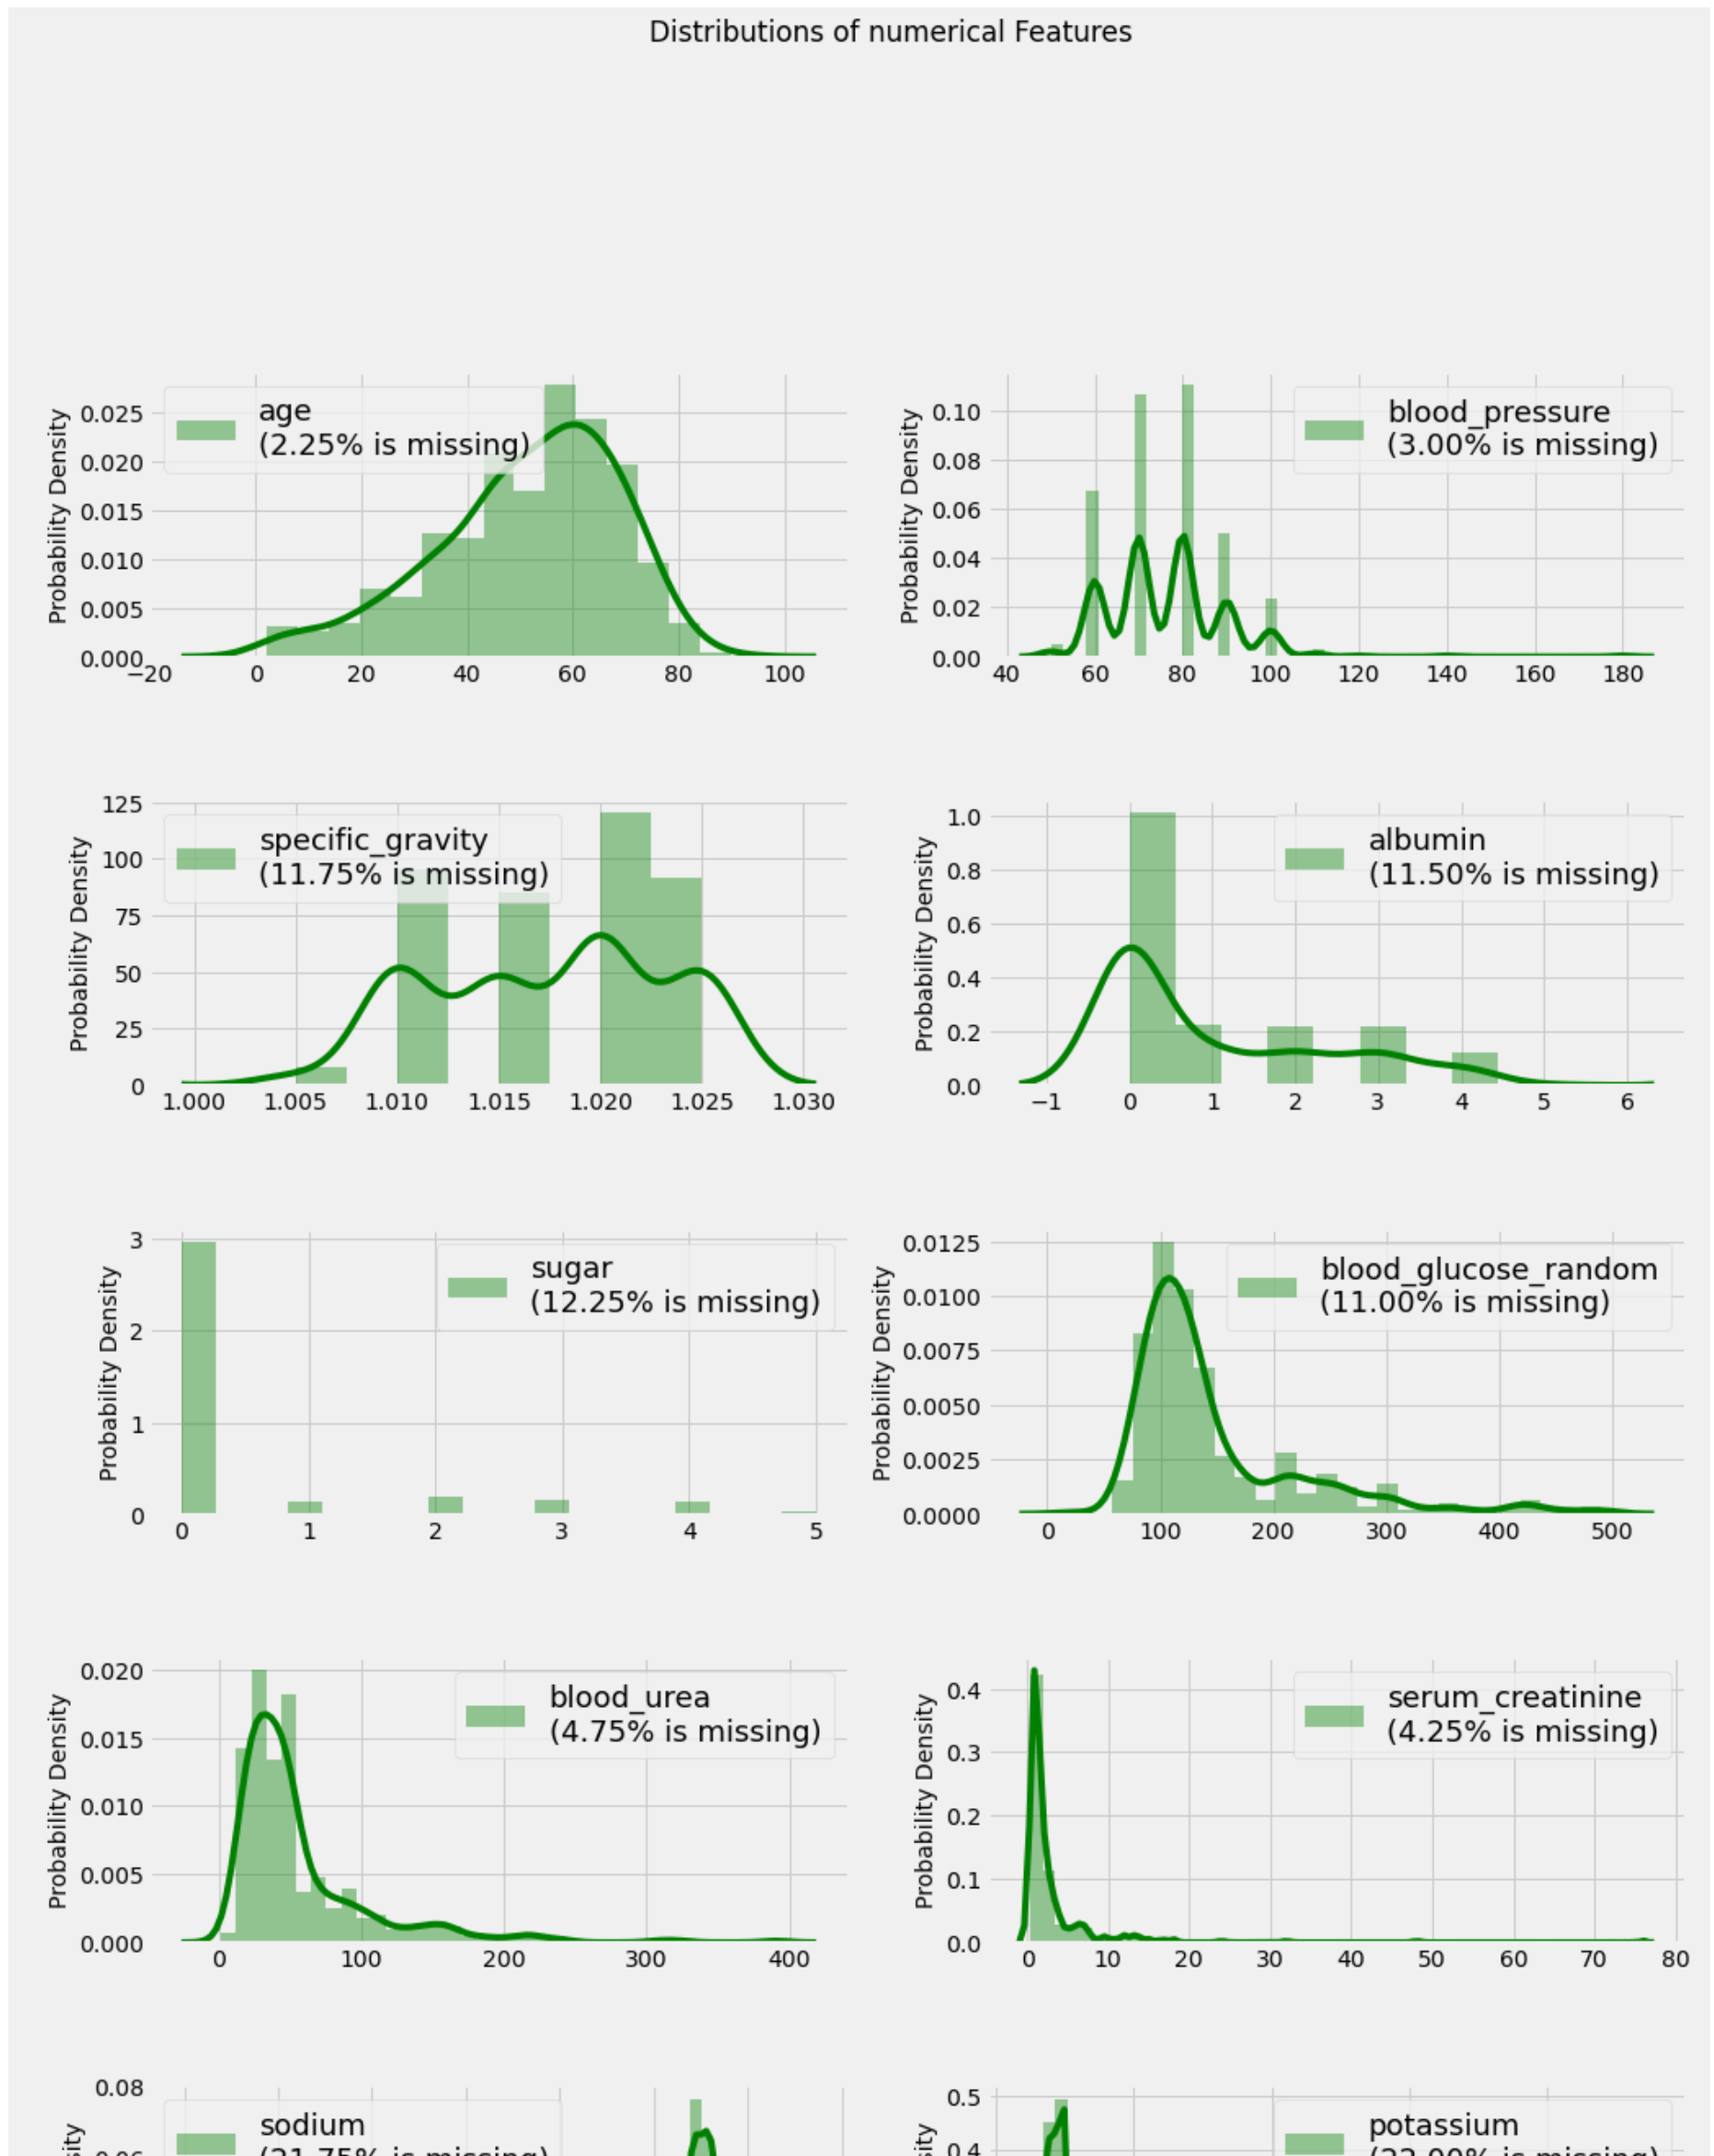
n_rows, n_cols = (7,2)

for index, column in enumerate(numeric):

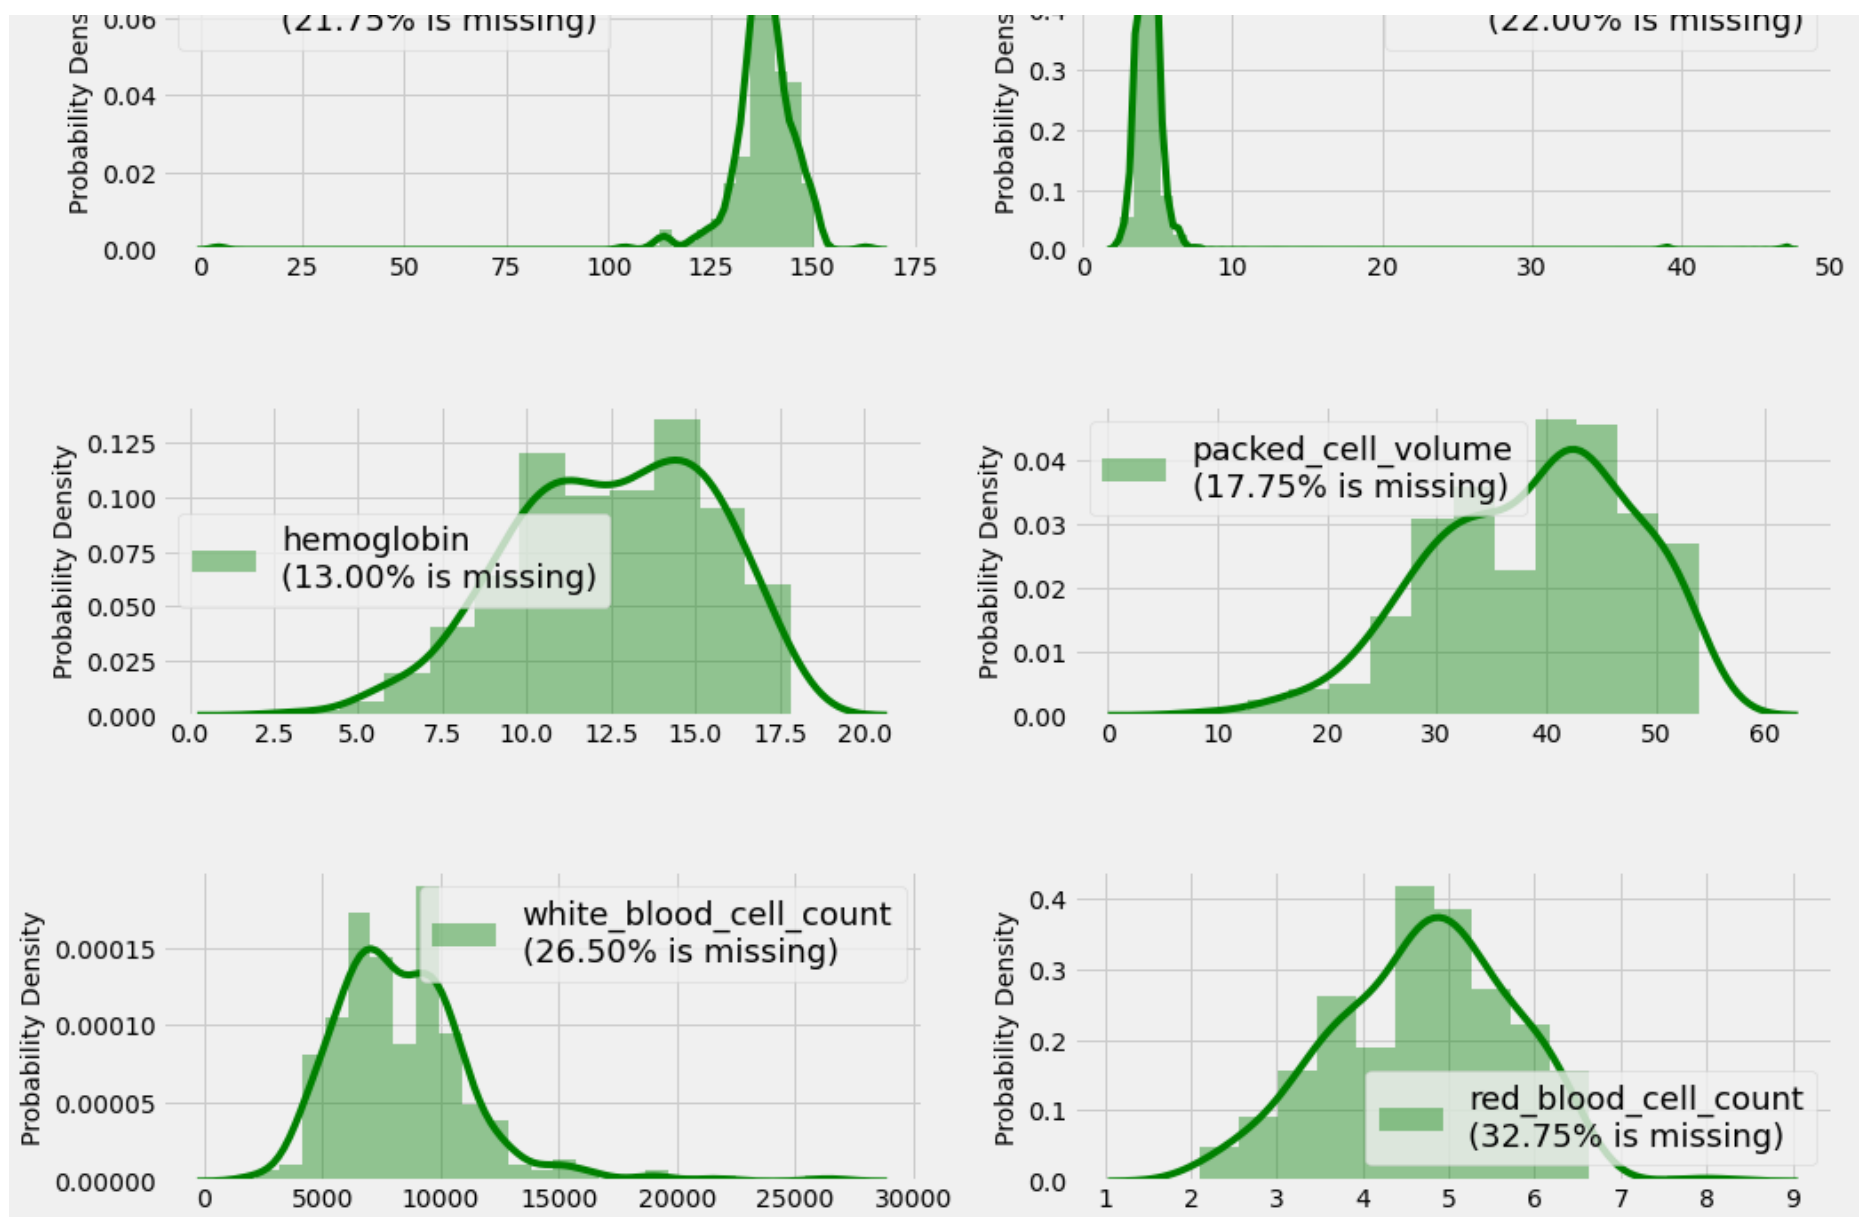
    i,j = (index // n_cols), (index % n_cols)
    miss_perc="%.2f"%(100*(1-(ckd_df[column].dropna().shape[0])/ckd_df.shape[0]))
    collabel=column+"\n({}% is missing)".format(miss_perc)
    fig=sns.distplot(ckd_df[column], color="green", label=collabel,
                    norm_hist=True, ax=axes[i,j], kde_kws={"lw":4})
    fig=fig.legend(loc='best', fontsize=18)

    axes[i,j].set_ylabel("Probability Density",fontsize='medium')
    axes[i,j].set_xlabel(None)

plt.show()
```







#### ***Numeric features:-***

While diving into detailed EDA we could see that some of the features represent good distributions, some are skewed by right & left and shows us positive and negative insights.

#### ***Checking distribution of the Categorical features***

```

In [15]: style.use('fivethirtyeight')

fig, axes = plt.subplots(nrows=5, ncols=2, figsize=(15,30))
fig.subplots_adjust(hspace=0.5)
fig.suptitle('Distributions of Categorical Features')

n_rows, n_cols = (5,2)

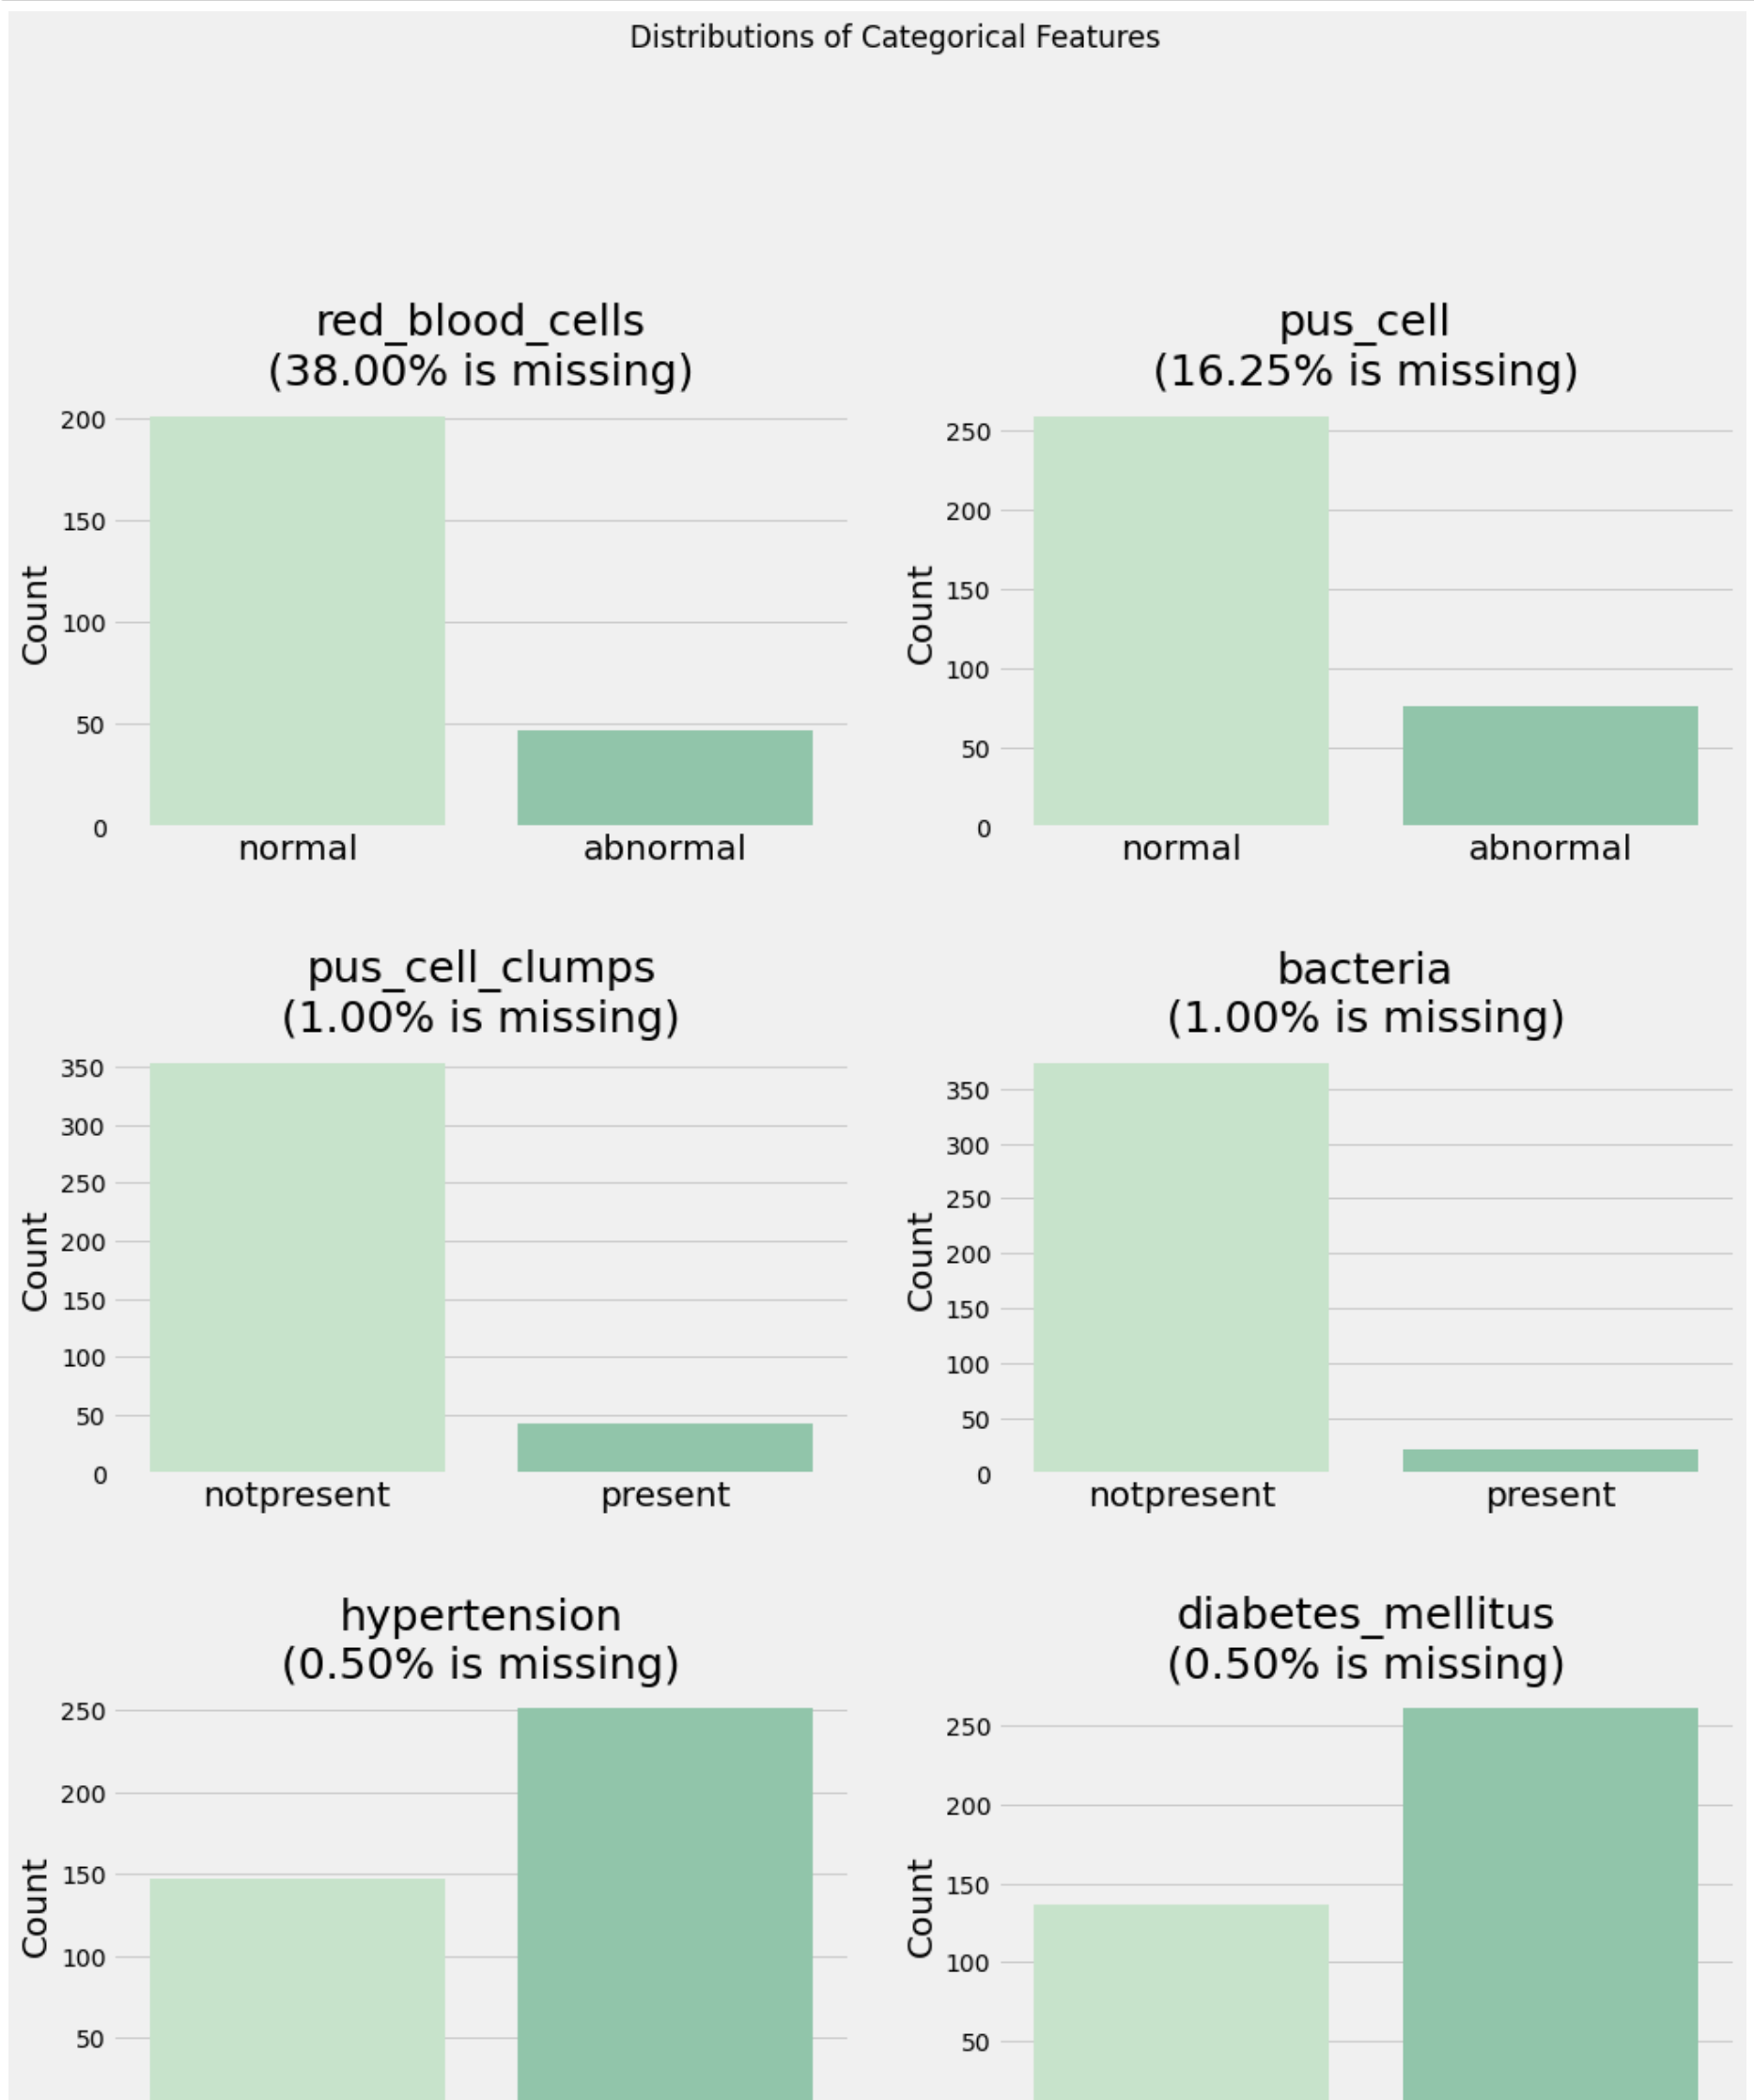
for index, column in enumerate(categoricals):


    i,j = index // n_cols, index % n_cols
    miss_perc="%.2f"%(100*(1-(ckd_df[column].dropna().shape[0])/ckd_df.shape[0]))
    collabel=column+"\n({}% is missing)".format(miss_perc)
    fig = sns.countplot(x=column, data=ckd_df,label=collabel,
                        palette=sns.cubehelix_palette(rot=-.4,light=0.85,hue=1), ax=axes[i,j])

    axes[i,j].set_title(collabel,fontsize=25)
    axes[i,j].set_xlabel(None)
    axes[i,j].set_ylabel("Count",fontsize=20)
    axes[i,j].set_xticklabels(axes[i,j].get_xticklabels(), Fontsize=20)

plt.show()

```





**Categorical features:-**

Bar graph is used to analysis categorical values. Since it's an binary classification dataset, we could see that Most of the unques features in categorical data we're 'good', 'poor', 'yes', 'no', 'normal', 'abnormal' and finally our classification label 'ckd' and 'notckd'.

```
In [16]: import matplotlib.ticker as ticker
style.use('fivethirtyeight')
# Some random data
ncount = 400

plt.figure(figsize=(15,8))
ax = sns.countplot(x="classification", data=ckd_df)
plt.title('Distribution of classification sata')
plt.xlabel('Type')

# Make twin axis
ax2=ax.twinx()

# Switch so count axis is on right, frequency on left
ax2.yaxis.tick_left()
ax.yaxis.tick_right()

# Also switch the labels over
ax.yaxis.set_label_position('right')
ax2.yaxis.set_label_position('left')

ax2.set_ylabel('Frequency [%]')

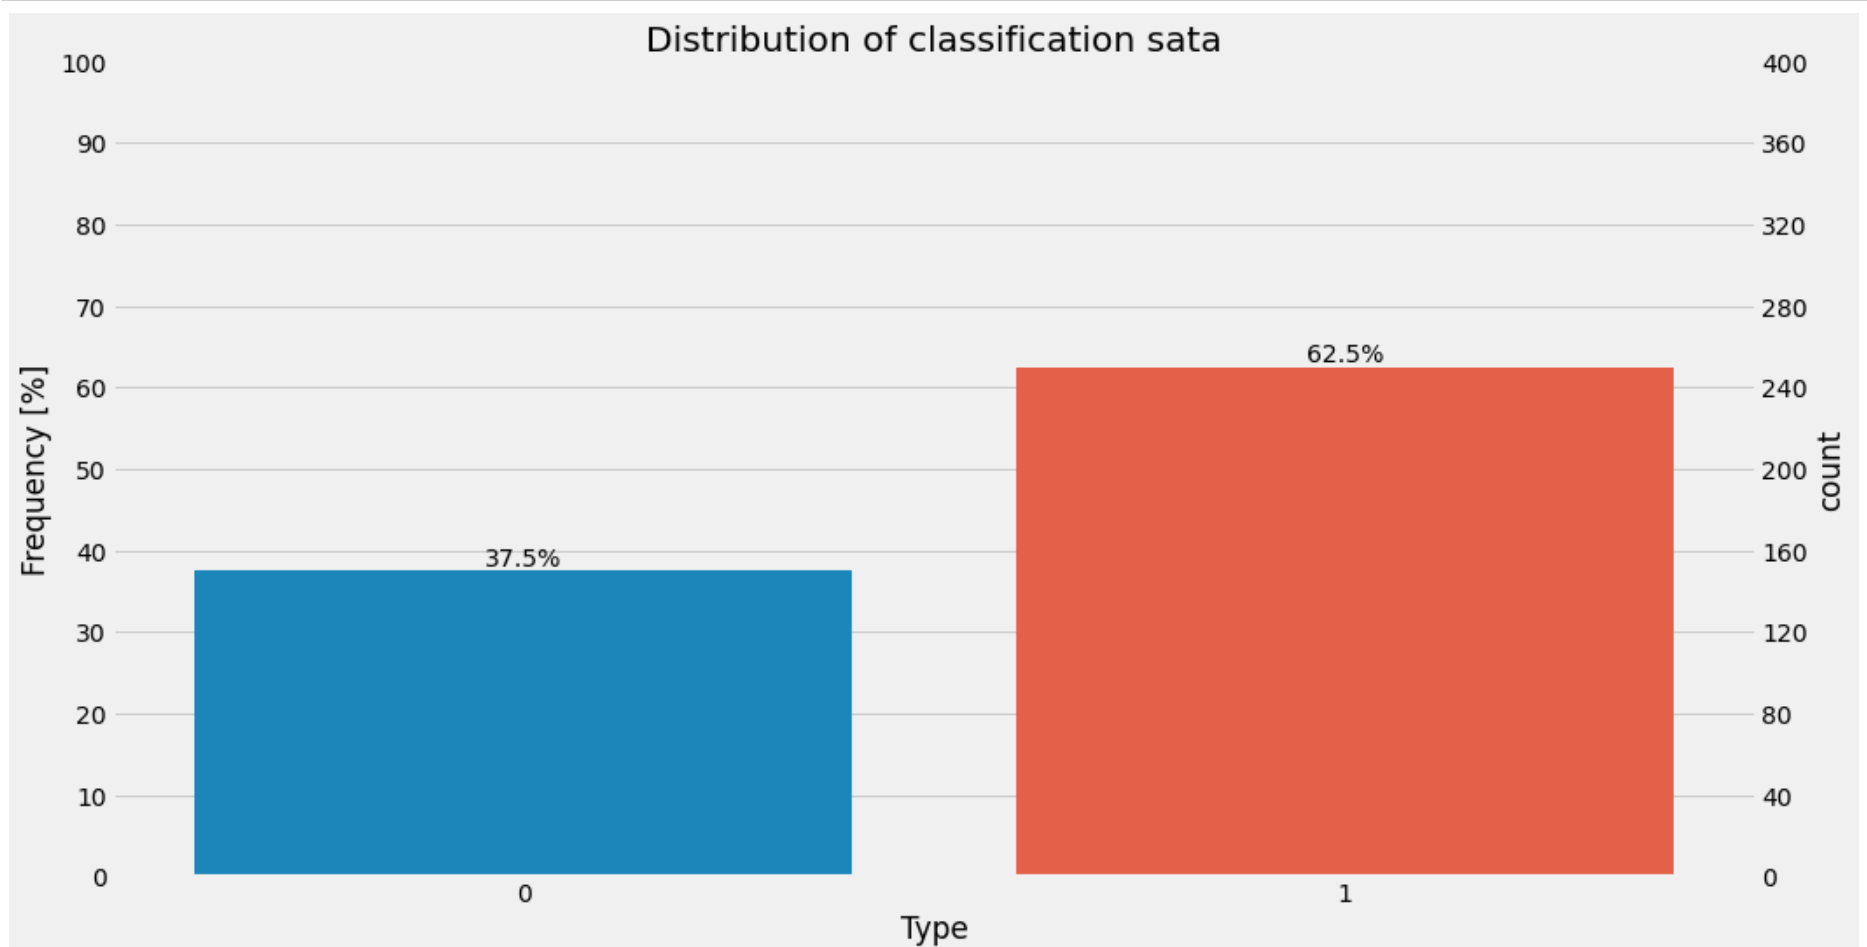
for p in ax.patches:
    x=p.get_bbox().get_points()[0,0]
    y=p.get_bbox().get_points()[1,1]
    ax.annotate('{:.1f}%'.format(100.*y/ncount), (x.mean(), y),
                ha='center', va='bottom') # set the alignment of the text

# Use a LinearLocator to ensure the correct number of ticks
ax.yaxis.set_major_locator(ticker.LinearLocator(11))

# Fix the frequency range to 0-100
ax2.set_ylim(0,100)
ax.set_ylim(0,ncount)

# And use a MultipleLocator to ensure a tick spacing of 10
ax2.yaxis.set_major_locator(ticker.MultipleLocator(10))

# Need to turn the grid on ax2 off, otherwise the gridlines end up on top of the bars
ax2.grid(None)
```



After required taking acting on Missing values using Sklearn package, Outliers were found by checking uniqueness values of each features.

Outliers removal:-

- Since we don't have higher level of data outliers.
- We took necessary outliers fix based on the domain knowledge by updating invalid data entry.
- Removing outliers will help to improve our data quality, which will be resulting in the model performance.


```
In [17]: for i in range(ckd_df.shape[0]):
    if ckd_df.iloc[i,24]=='ckd':
        ckd_df.iloc[i,24]='1'
    if ckd_df.iloc[i,24]=='notckd':
        ckd_df.iloc[i,24]='0'
```

## Missing Values

```
In [18]: style.use('fivethirtyeight')

d=((ckd_df.isnull().sum()/ckd_df.shape[0])).sort_values(ascending=False)
d.plot(kind='bar',
        color=sns.cubehelix_palette(start=2,
                                     rot=0.15,
                                     dark=0.15,
                                     light=0.95,
                                     reverse=True,
                                     n_colors=24),

        figsize=(20,10))
plt.title("\nProportions of Missing Values:\n",fontsize=40)
plt.show()
```



## Data cleaning

Data cleaning process were the key to projecting model quality and good accuracy score.

In CKD dataset we can clearly view the percentage of Missing values. Which is later handled using Sklearn missing value imputer.

Sklearn is one of the best library for imputation.

In [19]: ckd\_df.head(10).T

Out[19]:

	0	1	2	3	4	5	6	7	8	9
age	48	7	62	48	51	60	68	24	52	53
blood_pressure	80	50	80	70	80	90	70	NaN	100	90
specific_gravity	1.02	1.02	1.01	1.005	1.01	1.015	1.01	1.015	1.015	1.02
albumin	1	4	2	4	2	3	0	2	3	2
sugar	0	0	3	0	0	0	0	4	0	0
red_blood_cells	NaN	NaN	normal	normal	normal	NaN	NaN	normal	normal	abnormal
pus_cell	normal	normal	normal	abnormal	normal	NaN	normal	abnormal	abnormal	abnormal
pus_cell_clumps	notpresent	notpresent	notpresent	present	notpresent	notpresent	notpresent	notpresent	present	present
bacteria	notpresent	notpresent	notpresent	notpresent	notpresent	notpresent	notpresent	notpresent	notpresent	notpresent
blood_glucose_random	121	NaN	423	117	106	74	100	410	138	70
blood_urea	36	18	53	56	26	25	54	31	60	107
serum_creatinine	1.2	0.8	1.8	3.8	1.4	1.1	24	1.1	1.9	7.2
sodium	NaN	NaN	NaN	111	NaN	142	104	NaN	NaN	114
potassium	NaN	NaN	NaN	2.5	NaN	3.2	4	NaN	NaN	3.7
hemoglobin	15.4	11.3	9.6	11.2	11.6	12.2	12.4	12.4	10.8	9.5
packed_cell_volume	44	38	31	32	35	39	36	44	33	29
white_blood_cell_count	7800	6000	7500	6700	7300	7800	NaN	6900	9600	12100
red_blood_cell_count	5.2	NaN	NaN	3.9	4.6	4.4	NaN	5	4	3.7
hypertension	yes	no	no	yes	no	yes	no	no	yes	yes
diabetes_mellitus	yes	no	yes	no	no	yes	no	yes	yes	yes
coronary_artery_disease	no	no	no	no	no	no	no	no	no	no
appetite	good	good	poor	poor	good	good	good	good	good	poor
pedal_edema	no	no	no	yes	no	yes	no	yes	no	no
anemia	no	no	yes	yes	no	no	no	no	yes	yes
classification	1	1	1	1	1	1	1	1	1	1

One-Hot Encoding

```
In [20]: onehotdata=pd.get_dummies(ckd_df,drop_first=True,prefix_sep=': ')
onehotdata.head(13).T
```

Out[20]:

	0	1	2	3	4	5	6	7	8	9	10	11	12
age	48.00	7.00	62.00	48.000	51.00	60.000	68.00	24.000	52.000	53.00	50.00	63.00	68.000
blood_pressure	80.00	50.00	80.00	70.000	80.00	90.000	70.00	NaN	100.000	90.00	60.00	70.00	70.000
specific_gravity	1.02	1.02	1.01	1.005	1.01	1.015	1.01	1.015	1.015	1.02	1.01	1.01	1.015
albumin	1.00	4.00	2.00	4.000	2.00	3.000	0.00	2.000	3.000	2.00	2.00	3.00	3.000
sugar	0.00	0.00	3.00	0.000	0.00	0.000	0.00	4.000	0.000	0.00	4.00	0.00	1.000
blood_glucose_random	121.00	NaN	423.00	117.000	106.00	74.000	100.00	410.000	138.000	70.00	490.00	380.00	208.000
blood_urea	36.00	18.00	53.00	56.000	26.00	25.000	54.00	31.000	60.000	107.00	55.00	60.00	72.000
serum_creatinine	1.20	0.80	1.80	3.800	1.40	1.100	24.00	1.100	1.900	7.20	4.00	2.70	2.100
sodium	NaN	NaN	NaN	111.000	NaN	142.000	104.00	NaN	NaN	114.00	NaN	131.00	138.000
potassium	NaN	NaN	NaN	2.500	NaN	3.200	4.00	NaN	NaN	3.70	NaN	4.20	5.800
hemoglobin	15.40	11.30	9.60	11.200	11.60	12.200	12.40	12.400	10.800	9.50	9.40	10.80	9.700
packed_cell_volume	44.00	38.00	31.00	32.000	35.00	39.000	36.00	44.000	33.000	29.00	28.00	32.00	28.000
white_blood_cell_count	7800.00	6000.00	7500.00	6700.000	7300.00	7800.000	NaN	6900.000	9600.000	12100.00	NaN	4500.00	12200.000
red_blood_cell_count	5.20	NaN	NaN	3.900	4.60	4.400	NaN	5.000	4.000	3.70	NaN	3.80	3.400
classification	1.00	1.00	1.00	1.000	1.00	1.000	1.00	1.000	1.000	1.00	1.00	1.00	1.000
red_blood_cells: normal	0.00	0.00	1.00	1.000	1.00	0.000	0.00	1.000	1.000	0.00	0.00	0.00	0.000
pus_cell: normal	1.00	1.00	1.00	0.000	1.00	0.000	1.00	0.000	0.000	0.00	0.00	0.00	1.000
pus_cell_clumps: present	0.00	0.00	0.00	1.000	0.00	0.000	0.00	0.000	1.000	1.00	1.00	1.00	1.000
bacteria: present	0.00	0.00	0.00	0.000	0.00	0.000	0.00	0.000	0.000	0.00	0.00	0.00	0.000
hypertension: yes	1.00	0.00	0.00	1.000	0.00	1.000	0.00	0.000	1.000	1.00	1.00	1.00	1.000
diabetes_mellitus: yes	1.00	0.00	1.00	0.000	0.00	1.000	0.00	1.000	1.000	1.00	1.00	1.00	1.000
coronary_artery_disease: yes	0.00	0.00	0.00	0.000	0.00	0.000	0.00	0.000	0.000	0.00	0.00	0.00	1.000
appetite: poor	0.00	0.00	1.00	1.000	0.00	0.000	0.00	0.000	0.000	1.00	0.00	1.00	1.000
pedal_edema: yes	0.00	0.00	0.00	1.000	0.00	1.000	0.00	1.000	0.000	0.00	0.00	1.00	1.000
anemia: yes	0.00	0.00	1.00	1.000	0.00	0.000	0.00	0.000	1.000	1.00	1.00	0.00	0.000

```
In [21]: # define imputer
from sklearn.impute import KNNImputer

imputer = KNNImputer(n_neighbors=5, weights='uniform', metric='nan_euclidean')

impute_columns=list(set(onehotdata.columns)-set(["classification"]))
print(impute_columns)

['coronary_artery_disease: yes', 'pedal_edema: yes', 'albumin', 'blood_pressure', 'hemoglobin', 'sugar', 'white_blood_cell_count', 'bacteria: present', 'serum_creatinine', 'potassium', 'hypertension: yes', 'red_blood_cells: normal', 'specific_gravity', 'packed_cell_volume', 'red_blood_cell_count', 'sodium', 'diabetes_mellitus: yes', 'appetite: poor', 'pus_cell_clumps: present', 'blood_urea', 'blood_glucose_random', 'pus_cell: normal', 'age', 'anemia: yes']
```

```
In [22]: imputer.fit(onehotdata[impute_columns])
```

Out[22]: KNNImputer()

```
In [23]: X_trans=pd.DataFrame(imputer.transform(onehotdata[impute_columns]), columns=impute_columns)
```

In [24]: X\_trans.head(13).T

Out[24]:

	0	1	2	3	4	5	6	7	8	9	10	11	12
coronary_artery_disease: yes	0.00	0.00	0.00	0.000	0.00	0.000	0.00	0.000	0.000	0.00	0.00	0.00	1.000
pedal_edema: yes	0.00	0.00	0.00	1.000	0.00	1.000	0.00	1.000	0.000	0.00	0.00	1.00	1.000
albumin	1.00	4.00	2.00	4.000	2.00	3.000	0.00	2.000	3.000	2.00	2.00	3.00	3.000
blood_pressure	80.00	50.00	80.00	70.000	80.00	90.000	70.00	74.000	100.000	90.00	60.00	70.00	70.000
hemoglobin	15.40	11.30	9.60	11.200	11.60	12.200	12.40	12.400	10.800	9.50	9.40	10.80	9.700
sugar	0.00	0.00	3.00	0.000	0.00	0.000	0.00	4.000	0.000	0.00	4.00	0.00	1.000
white_blood_cell_count	7800.00	6000.00	7500.00	6700.000	7300.00	7800.000	10280.00	6900.000	9600.000	12100.00	9260.00	4500.00	12200.000
bacteria: present	0.00	0.00	0.00	0.000	0.00	0.000	0.00	0.000	0.000	0.00	0.00	0.00	0.000
serum_creatinine	1.20	0.80	1.80	3.800	1.40	1.100	24.00	1.100	1.900	7.20	4.00	2.70	2.100
potassium	4.20	3.92	4.20	2.500	3.98	3.200	4.00	4.200	4.960	3.70	4.56	4.20	5.800
hypertension: yes	1.00	0.00	0.00	1.000	0.00	1.000	0.00	0.000	1.000	1.00	1.00	1.00	1.000
red_blood_cells: normal	0.00	0.00	1.00	1.000	1.00	0.000	0.00	1.000	1.000	0.00	0.00	0.00	0.000
specific_gravity	1.02	1.02	1.01	1.005	1.01	1.015	1.01	1.015	1.015	1.02	1.01	1.01	1.015
packed_cell_volume	44.00	38.00	31.00	32.000	35.00	39.000	36.00	44.000	33.000	29.00	28.00	32.00	28.000
red_blood_cell_count	5.20	4.96	3.80	3.900	4.60	4.400	4.64	5.000	4.000	3.70	4.92	3.80	3.400
sodium	137.60	136.80	133.80	111.000	138.40	142.000	104.00	133.800	134.000	114.00	139.00	131.00	138.000
diabetes_mellitus: yes	1.00	0.00	1.00	0.000	0.00	1.000	0.00	1.000	1.000	1.00	1.00	1.00	1.000
appetite: poor	0.00	0.00	1.00	1.000	0.00	0.000	0.00	0.000	0.000	1.00	0.00	1.00	1.000
pus_cell_clumps: present	0.00	0.00	0.00	1.000	0.00	0.000	0.00	0.000	1.000	1.00	1.00	1.00	1.000
blood_urea	36.00	18.00	53.00	56.000	26.00	25.000	54.00	31.000	60.000	107.00	55.00	60.00	72.000
blood_glucose_random	121.00	113.00	423.00	117.000	106.00	74.000	100.00	410.000	138.000	70.00	490.00	380.00	208.000
pus_cell: normal	1.00	1.00	1.00	0.000	1.00	0.000	1.00	0.000	0.000	0.00	0.00	0.00	1.000
age	48.00	7.00	62.00	48.000	51.00	60.000	68.00	24.000	52.000	53.00	50.00	63.00	68.000
anemia: yes	0.00	0.00	1.00	1.000	0.00	0.000	0.00	0.000	1.000	1.00	1.00	0.00	0.000

In [25]: X\_trans # final dataset

Out[25]:

	coronary_artery_disease: yes	pedal_edema: yes	albumin	blood_pressure	hemoglobin	sugar	white_blood_cell_count	bacteria: present	serum_creatinine	potass
0	0.0	0.0	1.0	80.0	15.4	0.0	7800.0	0.0	1.2	
1	0.0	0.0	4.0	50.0	11.3	0.0	6000.0	0.0	0.8	
2	0.0	0.0	2.0	80.0	9.6	3.0	7500.0	0.0	1.8	
3	0.0	1.0	4.0	70.0	11.2	0.0	6700.0	0.0	3.8	
4	0.0	0.0	2.0	80.0	11.6	0.0	7300.0	0.0	1.4	
...	...	...	...	...	...	...	...	...	...	
395	0.0	0.0	0.0	80.0	15.7	0.0	6700.0	0.0	0.5	
396	0.0	0.0	0.0	70.0	16.5	0.0	7800.0	0.0	1.2	
397	0.0	0.0	0.0	80.0	15.8	0.0	6600.0	0.0	0.6	
398	0.0	0.0	0.0	60.0	14.2	0.0	7200.0	0.0	1.0	
399	0.0	0.0	0.0	80.0	15.8	0.0	6800.0	0.0	1.1	

400 rows × 24 columns

Modelling



```
In [26]: X=X_trans
y=ckd_df["classification"]
```

```
X_prod=X_trans
print(X.shape)
print(y.shape)
print(X_prod.shape)
```

```
(400, 24)
(400,)
(400, 24)
```

### ***Predictive Models with hyperparameter tuning Section***

```
In [27]: from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix

from sklearn.model_selection import GridSearchCV
```

```
In [28]: def display_confusion_matrix(y_test,y_pred):

    cm = confusion_matrix(y_test, y_pred_lr)
    group_names = ["True Neg", "False Pos", "False Neg", "True Pos"]
    group_counts = ["{0:0.0f}".format(value) for value in cm.flatten()]
    group_percentages = ["{0:.2%}".format(value) for value in cm.flatten()/np.sum(cm)]

    labels = [f"{v1}\n{v2}\n{v3}" for v1, v2, v3 in zip(group_names,group_counts,group_percentages)]
    labels = np.asarray(labels).reshape(2,2)

    sns.heatmap(cm, annot=labels, fmt="", cmap="Blues")
    print(classification_report(y_test, y_pred))
```

```
In [29]: def plot_roc_curve(fpr, tpr):
    plt.plot(fpr, tpr, label='ROC')
    plt.plot([0, 1], linestyle='--')
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.legend()
    plt.show()
```

```
In [30]: ##Split train and test
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20, random_state = 4658)

print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(320, 24)
(80, 24)
(320,)
(80,)
```

### ***StandardScaler data with PCA implementation***

```
#performing preprocessing part from sklearn.preprocessing import StandardScaler sc = StandardScaler()
```

```
X_train = sc.fit_transform(X_train) X_test = sc.transform(X_test)
```

```
#Applying PCA function on training and testing set of X component from sklearn.decomposition import PCA
```

```
pca = PCA(n_components = 2)
```

```
X_train = pca.fit_transform(X_train) X_test = pca.transform(X_test)
```

```
explained_variance = pca.explained_variance_ratio_ explained_variance
```

### ***Logistic Regression Hyper parameter tuning***

```
In [31]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

c_space = np.logspace(-5, 8, 15)
param_grid = {'C': c_space}

start_lr = time.time()
lr = GridSearchCV(LogisticRegression(),
                  param_grid,
                  cv = 5)
lr.fit(X_train, y_train)
end_lr = time.time()
final_lr = end_lr - start_lr
final_lr = round(final_lr,3)
final_lr


# Print the tuned parameters and score
print("Tuned Logistic Regression Parameters: {}".format(lr.best_params_))
print("Best score is {}".format(lr.best_score_))
print("Best estimator is {} \n\n".format(lr.best_estimator_))

y_pred_lr = lr.predict(X_test)
display_confusion_matrix(y_test, y_pred_lr)
accuracy_lr=lr.best_score_
print("\nAccuracy of Logistic Regression is :", accuracy_lr)
print("Computation time {} - Sec".format(final_lr))
```

```
Tuned Logistic Regression Parameters: {'C': 2275.845926074791}
Best score is 0.9
Best estimator is LogisticRegression(C=2275.845926074791)
```

	precision	recall	f1-score	support
0	0.84	0.84	0.84	25
1	0.93	0.93	0.93	55
accuracy			0.90	80
macro avg	0.88	0.88	0.88	80
weighted avg	0.90	0.90	0.90	80


```
Accuracy of Logistic Regression is : 0.9
Computation time 3.003 - Sec
```



```
In [32]: auc = roc_auc_score(y_test, y_pred_lr)
print('AUC: %.2f' % auc)

fpr, tpr, thresholds = roc_curve(y_test, y_pred_lr)
plot_roc_curve(fpr, tpr)
```

AUC: 0.88



**Decision Tree Hyper parameter tuning**

```
In [33]: from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import RandomizedSearchCV

hyperparam_combs = {
    'max_depth': [4, 6, 8, 10, 12],
    'criterion': ['gini', 'entropy'],
    'min_samples_split': [2, 10, 20, 30, 40],
    'max_features': [0.2, 0.4, 0.6, 0.8, 1],
    'max_leaf_nodes': [8, 16, 32, 64, 128],
    'class_weight': [{0: 1, 1: 1}, {0: 1, 1: 2}, {0: 1, 1: 3}, {0: 1, 1: 4}, {0: 1, 1: 5}]
}

start_dt = time.time()
clf = RandomizedSearchCV(DecisionTreeClassifier(),
                        hyperparam_combs,
                        scoring='f1',
                        random_state=1,
                        n_iter=20)

dt_model = clf.fit(X_train, y_train)
end_dt = time.time()
final_dt = end_dt - start_dt
final_dt = round(final_dt,3)
final_dt


# Print the tuned parameters and score
print("Tuned Decision Tree Parameters: {}".format(dt_model.best_params_))
print("Best score is {}".format(dt_model.best_score_))
print("Best estimator is {}".format(dt_model.best_estimator_))

y_pred_dt = dt_model.predict(X_test)
display_confusion_matrix(y_test, y_pred_dt)
accuracy_dt=dt_model.best_score_
print("Accuracy of Decision Tree is :", accuracy_dt)
print("Computation time {} - Sec".format(final_dt))
```

```
Tuned Decision Tree Parameters: {'min_samples_split': 2, 'max_leaf_nodes': 128, 'max_features': 0.2, 'max_depth': 10, 'criterion': 'gini', 'class_weight': {0: 1, 1: 3}}
Best score is 0.9687584678257564
Best estimator is DecisionTreeClassifier(class_weight={0: 1, 1: 3}, max_depth=10,
                                         max_features=0.2, max_leaf_nodes=128)
```

	precision	recall	f1-score	support
0	0.89	0.96	0.92	25
1	0.98	0.95	0.96	55
accuracy			0.95	80
macro avg	0.94	0.95	0.94	80
weighted avg	0.95	0.95	0.95	80


```
Accuracy of Decision Tree is : 0.9687584678257564
Computation time 0.609 - Sec
```



```
In [34]: auc = roc_auc_score(y_test, y_pred_dt)
print('AUC: %.2f' % auc)

fpr, tpr, thresholds = roc_curve(y_test, y_pred_dt)
plot_roc_curve(fpr, tpr)
```

AUC: 0.95



***Random Forest Hyper parameter tuning***

```
In [35]: from sklearn.ensemble import RandomForestClassifier
```

```
param_grid = {"n_estimators": np.arange(2, 300, 2),
              "max_depth": np.arange(1, 28, 1),
              "min_samples_split": np.arange(1, 150, 1),
              "min_samples_leaf": np.arange(1, 60, 1),
              "max_leaf_nodes": np.arange(2, 60, 1),
              "min_weight_fraction_leaf": np.arange(0.1, 0.4, 0.1)}
```

```
start_rf = time.time()
rf = RandomizedSearchCV(RandomForestClassifier(),
                        param_grid,
                        scoring='f1',
                        random_state=4658,
                        n_iter=20)
```

```
rf_model = rf.fit(X_train, y_train)
end_rf = time.time()
final_rf = end_rf - start_rf
final_rf = round(final_rf, 3)
final_rf
```

```
# Print the tuned parameters and score
print("Tuned Random Tree Parameters: {}".format(rf_model.best_params_))
print("Best score is {}".format(rf_model.best_score_))
print("Best estimator is {}".format(rf_model.best_estimator_))
```

```
y_pred_rf = rf_model.predict(X_test)
display_confusion_matrix(y_test, y_pred_rf)
accuracy_rf = rf_model.best_score_
print("Accuracy of Random Forests model is :", accuracy_rf)
print("Computation time {} - Sec".format(final_rf))
```

```
Tuned Random Tree Parameters: {'n_estimators': 240, 'min_weight_fraction_leaf': 0.1, 'min_samples_split': 52,
'min_samples_leaf': 38, 'max_leaf_nodes': 39, 'max_depth': 26}
```

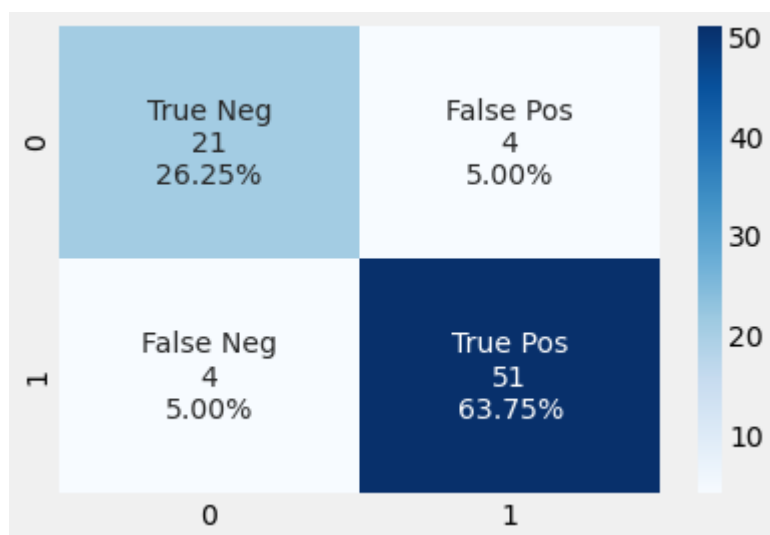
```
Best score is 0.9740925740925741
```

```
Best estimator is RandomForestClassifier(max_depth=26, max_leaf_nodes=39, min_samples_leaf=38,
min_samples_split=52, min_weight_fraction_leaf=0.1,
n_estimators=240)
```

	precision	recall	f1-score	support
0	0.96	1.00	0.98	25
1	1.00	0.98	0.99	55
accuracy			0.99	80
macro avg	0.98	0.99	0.99	80
weighted avg	0.99	0.99	0.99	80


```
Accuracy of Random Forests model is : 0.9740925740925741
```

```
Computation time 22.088 - Sec
```



```
In [36]: auc = roc_auc_score(y_test, y_pred_rf)
print('AUC: %.2f' % auc)

fpr, tpr, thresholds = roc_curve(y_test, y_pred_rf)
plot_roc_curve(fpr, tpr)
```



```
In [37]: from sklearn.svm import SVC

# defining parameter range
param_grid = {'C': [0.1, 1, 10, 100, 1000],
              'gamma': [1, 0.1, 0.01, 0.001, 0.0001]}

start_svm = time.time()
svm = GridSearchCV(SVC(), param_grid, refit = True, verbose = 3)
# fitting the model for grid search
svm.fit(X_train, y_train)

end_svm = time.time()
final_svm = end_svm - start_svm
final_svm = round(final_svm,3)
final_svm

# Print the tuned parameters and score
print("Tuned Support Vector Machine Parameters: {}".format(svm.best_params_))
print("Best score is {}".format(svm.best_score_))
print("Best estimator is {}".format(svm.best_estimator_))
```


---

```
[CV] ..... C=0.1, gamma=1, score=0.609, total= 0.0s
[CV] C=0.1, gamma=1 .....
[CV] ..... C=0.1, gamma=1, score=0.609, total= 0.0s
[CV] C=0.1, gamma=1 .....
[CV] ..... C=0.1, gamma=1, score=0.609, total= 0.0s
[CV] C=0.1, gamma=0.1 .....
[CV] ..... C=0.1, gamma=0.1, score=0.609, total= 0.0s
[CV] C=0.1, gamma=0.1 .....
[CV] ..... C=0.1, gamma=0.1, score=0.609, total= 0.0s
[CV] C=0.1, gamma=0.1 .....
[CV] ..... C=0.1, gamma=0.1, score=0.609, total= 0.0s
[CV] C=0.1, gamma=0.1 .....
[CV] ..... C=0.1, gamma=0.1, score=0.609, total= 0.0s
[CV] C=0.1, gamma=0.1 .....
[CV] ..... C=0.1, gamma=0.1, score=0.609, total= 0.0s
[CV] C=0.1, gamma=0.1 .....
[CV] ..... C=0.1, gamma=0.1, score=0.609, total= 0.0s
[CV] C=0.1, gamma=0.01 .....
[CV] ..... C=0.1, gamma=0.01, score=0.609, total= 0.0s
[CV] C=0.1, gamma=0.01 .....
[CV] ..... C=0.1, gamma=0.01, score=0.609, total= 0.0s
[CV] C=0.1, gamma=0.01 .....
```

```
In [38]: y_pred_svm = svm.predict(X_test)
display_confusion_matrix(y_test, y_pred_svm)
accuracy_svm=svm.best_score_
print("Accuracy of Support Vector Machine is :", accuracy_svm)
print("Computation time {} - Sec".format(final_svm))
```

	precision	recall	f1-score	support
0	0.83	0.60	0.70	25
1	0.84	0.95	0.89	55
accuracy			0.84	80
macro avg	0.84	0.77	0.79	80
weighted avg	0.84	0.84	0.83	80

Accuracy of Support Vector Machine is : 0.75625  
Computation time 1.222 - Sec



```
In [39]: auc = roc_auc_score(y_test, y_pred_svm)
print('AUC: %.2f' % auc)

fpr, tpr, thresholds = roc_curve(y_test, y_pred_svm)
plot_roc_curve(fpr, tpr)
```

AUC: 0.77



**Artificial neural network**



```
In [40]: from sklearn.neural_network import MLPClassifier
```

```
# defining parameter range
param_grid = {
    'max_iter': [1000],
}

start_mlp = time.time()
mlp = GridSearchCV(MLPClassifier(), param_grid, refit = True, verbose = 3)
# fitting the model for grid search
mlp.fit(X_train, y_train.values.ravel())

end_mlp = time.time()
final_mlp = end_mlp - start_mlp
final_mlp = round(final_mlp,3)
final_mlp

# Print the tuned parameters and score
print("Tuned Artificial neural network Parameters: {}".format(mlp.best_params_))
print("Best score is {}".format(mlp.best_score_))
print("Best estimator is {}".format(mlp.best_estimator_))
```

```
Fitting 5 folds for each of 1 candidates, totalling 5 fits
[CV] max_iter=1000 .....
[CV] ..... max_iter=1000, score=0.703, total= 0.2s
[CV] max_iter=1000 .....

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 0.2s remaining: 0.0s

[CV] ..... max_iter=1000, score=0.859, total= 0.2s
[CV] max_iter=1000 .....
[CV] ..... max_iter=1000, score=0.875, total= 0.2s
[CV] max_iter=1000 .....

[Parallel(n_jobs=1)]: Done 2 out of 2 | elapsed: 0.4s remaining: 0.0s

[CV] ..... max_iter=1000, score=0.734, total= 0.2s
[CV] max_iter=1000 .....
[CV] ..... max_iter=1000, score=0.781, total= 0.4s


[Parallel(n_jobs=1)]: Done 5 out of 5 | elapsed: 1.2s finished

Tuned Artificial neural network Parameters: {'max_iter': 1000}
Best score is 0.790625
Best estimator is MLPClassifier(max_iter=1000)
```

```
In [41]: y_pred_mlp = mlp.predict(X_test)
display_confusion_matrix(y_test, y_pred_mlp)
accuracy_mlp=mlp.best_score_
print("Accuracy of Artificial neural network is :", accuracy_mlp)
print("Computation time {} - Sec".format(final_mlp))
```

	precision	recall	f1-score	support
0	0.77	0.96	0.86	25
1	0.98	0.87	0.92	55
accuracy			0.90	80
macro avg	0.88	0.92	0.89	80
weighted avg	0.92	0.90	0.90	80


```
Accuracy of Artificial neural network is : 0.790625
Computation time 1.735 - Sec
```



```
In [42]: auc = roc_auc_score(y_test, y_pred_mlp)
print('AUC: %.2f' % auc)

fpr, tpr, thresholds = roc_curve(y_test, y_pred_mlp)
plot_roc_curve(fpr, tpr)
```

AUC: 0.92



***k* Nearest Neighbors**

```
In [43]: from sklearn.neighbors import KNeighborsClassifier

model = KNeighborsClassifier(n_jobs=-1)
params = {'n_neighbors':[5,6,7,8,9,10],
          'leaf_size':[1,2,3,5],
          'weights':['uniform', 'distance'],
          'algorithm':['auto', 'ball_tree','kd_tree','brute'],
          'n_jobs':[-1]}

start_knn = time.time()
modell = GridSearchCV(model, param_grid=params, n_jobs=1)
modell.fit(X_train, y_train)
end_knn = time.time()
final_knn = end_knn - start_knn
final_knn = round(final_knn,3)
final_knn


# Print the tuned parameters and score
print("Tuned Logistic Regression Parameters: {}".format(modell.best_params_))
print("Best score is {}".format(modell.best_score_))
print("Best estimator is {} \n\n".format(modell.best_estimator_))

y_pred_knn = modell.predict(X_test)
display_confusion_matrix(y_test, y_pred_knn)
accuracy_knn=accuracy_score(y_test, y_pred_knn)
print("\nAccuracy of k Nearest Neighbors is \t:", accuracy_knn)
print("Best Accuracy of k Nearest Neighbors is : {}".format(modell.best_score_))
print("Computation time {} - Sec".format(final_knn))
```

Tuned Logistic Regression Parameters: {'algorithm': 'auto', 'leaf\_size': 1, 'n\_jobs': -1, 'n\_neighbors': 7, 'weights': 'distance'}  
Best score is 0.68125  
Best estimator is KNeighborsClassifier(leaf\_size=1, n\_jobs=-1, n\_neighbors=7, weights='distance')

	precision	recall	f1-score	support
0	0.44	0.64	0.52	25
1	0.80	0.64	0.71	55
accuracy			0.64	80
macro avg	0.62	0.64	0.62	80
weighted avg	0.69	0.64	0.65	80


Accuracy of k Nearest Neighbors is : 0.6375  
Best Accuracy of k Nearest Neighbors is : 0.68125  
Computation time 16.407 - Sec



```
In [44]: auc = roc_auc_score(y_test, y_pred_knn)
print('AUC: %.2f' % auc)

fpr, tpr, thresholds = roc_curve(y_test, y_pred_knn)
plot_roc_curve(fpr, tpr)
```

AUC: 0.64



### Naive Bayes (NB) Classifier

```
In [45]: from sklearn.naive_bayes import GaussianNB

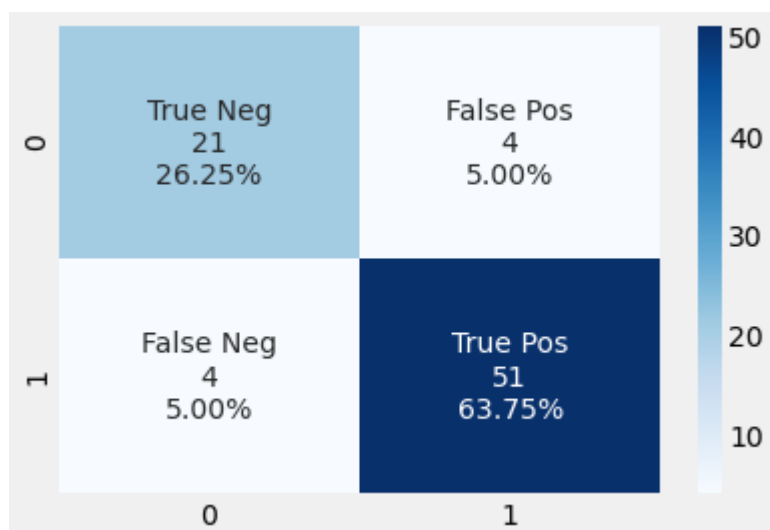
start_nb = time.time()
gnb = GaussianNB()
gnb.fit(X_train, y_train)
nb_model = gnb.fit(X_train, y_train)
end_nb = time.time()
final_nb = end_nb - start_nb
final_nb = round(final_nb,3)
final_nb

y_pred_nb = gnb.predict(X_test)

display_confusion_matrix(y_test, y_pred_nb)
accuracy_nb=accuracy_score(y_test, y_pred_nb)
print("Gaussian Naive Bayes model accuracy :", accuracy_nb)
print("Computation time {} - Sec".format(final_nb))
```

	precision	recall	f1-score	support
0	0.83	1.00	0.91	25
1	1.00	0.91	0.95	55
accuracy			0.94	80
macro avg	0.92	0.95	0.93	80
weighted avg	0.95	0.94	0.94	80


Gaussian Naive Bayes model accuracy : 0.9375  
Computation time 0.005 - Sec



```
In [46]: auc = roc_auc_score(y_test, y_pred_nb)
print('AUC: %.2f' % auc)

fpr, tpr, thresholds = roc_curve(y_test, y_pred_nb)
plot_roc_curve(fpr, tpr)
```

AUC: 0.95



```
In [47]: models = pd.DataFrame({
    'Model': ['Logistic Regression', 'Naive Bayes', 'Decision Tree', 'Random Forest', 'Support Vector Machines',
              'K - Nearest Neighbors', 'Artificial neural network'],
    'Score': [accuracy_lr, accuracy_nb , accuracy_dt, accuracy_rf, svm.best_score_, model1.best_score_, mlp.best_score_],
    'Rank': range(7)}).sort_values(by='Score', ascending=False)
```


Out[47]:

	Model	Score
3	Random Forest	0.974093
2	Decision Tree	0.968758
1	Naive Bayes	0.937500
0	Logistic Regression	0.900000
6	Artificial neural network	0.790625
4	Support Vector Machines	0.756250
5	K - Nearest Neighbors	0.681250

```
In [48]: import plotly.express as px
import plotly.graph_objects as go

fig = px.bar(models, x='Model', y='Score', color="Model", title="Model Comparison")
fig.show()
```

Model Comparison




```
In [49]: model = models["Model"]

fig = go.Figure()
fig.add_trace(go.Bar(x=model,
                    y=models["Score"],
                    name='Without dimensionality reduction',
                    marker_color='rgb(55, 83, 109)'
                    ))
fig.add_trace(go.Bar(x=model,
                    y=[0.978125,1.000000, 0.987277, 0.987277, 0.987500, 0.984375, 0.981250],
                    name='With PCA',
                    marker_color='rgb(26, 118, 255)'
                    ))

fig.update_layout(
    title='US Export of Plastic Scrap',
    xaxis_tickfont_size=10,
    yaxis=dict(
        title='USD (millions)',
        titlefont_size=16,
        tickfont_size=14,
    ),
    legend=dict(
        x=0,
        y=1.0,
        bgcolor='rgba(255, 255, 255, 0)',
        bordercolor='rgba(255, 255, 255, 0)'
    ),
    barmode='group',
    bargap=0.15, # gap between bars of adjacent location coordinates.
    bargroupgap=0.1 # gap between bars of the same location coordinate.
)
fig.show()
```

US Export of Plastic Scrap



Insights,

24 data recording information of 400 people such as age, blood pressure, density, diabetes were used as attributes. Clinical records were examined to determine whether chronic kidney disease was present or not, and provided a high accuracy rate with machine learning methods.

Chronic kidney disease is a disease that hinders the normal functions of the kidney and damages the kidneys. It is one of the common diseases in the world and the prediction of the disease is one of the basic issues in medical diagnosis. Chronic kidney disease is one of the leading causes of death worldwide. Early detection of this disease is very important in terms of health and treatment costs. Many machine learning algorithms have been used in the literature to predict the disease.

Here, seven different classifiers were utilized in determining the targeted chronic kidney disease and the best performing classifier was tried to be found. These algorithms were compared on the basis of accuracy, sensitivity, sensitivity, recall and f1 score. When the results were evaluated with the data used in this study, it was seen that the **Random Forest and SVM** method (with an accuracy of 99.16%) performed better than other classification algorithms.

Here by concluding the **Naive Bayes** is normally giving good accuracy score with taken dataset and notably **Naive Bayes** is overfitting after implementing dimensionality reduction with PCA (2 components)

